

RAJALAKSHMI ENGINEERING COLLEGE
RAJALAKSHMI NAGAR, THANDALAM – 602 105



RAJALAKSHMI
ENGINEERING COLLEGE
An AUTONOMOUS Institution
Affiliated to ANNA UNIVERSITY, Chennai

**CS19443 DATABASE MANAGEMENT
SYSTEMS LAB**

LABORATORY RECORD NOTE BOOK

Name : **RAKHUL PRAKASH S B**

Year / Branch / Section : **II / B.E CSE / D**

University Register No. : **2116220701216**

College Roll No. : **220701216**

Semester : **IV**

Academic Year : **2023-2024**

INDEX

Name : _____ Branch : _____ Sec : _____ Roll No : _____

[illegible]

Definition of a Relational Database

A relational database is a collection of relations or two-dimensional tables.

Terminologies Used in a Relational Database

1. A single **ROW** or table representing all data required for a particular employee. Each row should be identified by a primary key which allows no duplicate rows.
2. A **COLUMN** or attribute containing the employee number which identifies a unique employee. Here Employee number is designated as a primary key ,must contain a value and must be unique.
3. A column may contain foreign key. Here Dept_ID is a foreign key in employee table and it is a primary key in Department table.
4. A Field can be found at the intersection of a row and column. There can be only one value in it. Also it may have no value. This is called a null value.

EMP ID	FIRST NAME	LAST NAME	EMAIL
100	King	Steven	Sking
101	John	Smith	Jsmith
102	Neena	Bai	Neenba
103	Eex	De Haan	Ldehaan

Relational Database Properties

A relational database :

- Can be accessed and modified by executing structured query language (SQL) statements.
- Contains a collection of tables with no physical pointers.
- Uses a set of operators

Relational Database Management Systems

RDBMS refers to a relational database plus supporting software for managing users and processing SQL queries, performing backups/restores and associated tasks. (Relational Database Management System) Software for storing data using SQL (structured query language). A relational database uses SQL to store data in a series of tables that not only record existing relationships between data items, but which also permit the data to be joined in new relationships. SQL (pronounced 'sequel') is based on a system of algebra developed by E F Codd, an IBM scientist who first defined the relational model in 1970. Relational databases are optimized for storing transactional data, and the majority of modern business software applications therefore use an RDBMS as their data store. The leading RDBMS vendors are Oracle, IBM and Microsoft.

The first commercial RDBMS was the Multics Relational Data Store, first sold in 1978. INGRES, Oracle, Sybase, Inc., Microsoft Access, and Microsoft SQL Server are well-known database products and companies. Others include PostgreSQL, SQL/DS, and RDB. A relational database management system (RDBMS) is a program that lets you create, update, and administer a relational database. Most commercial RDBMS's use the Structured Query Language (SQL) to access the database, although SQL was invented after the development of the relational model and is not necessary for its use.

The leading RDBMS products are Oracle, IBM's DB2 and Microsoft's SQL Server. Despite repeated challenges by competing technologies, as well as the claim by some experts that no current RDBMS has fully implemented relational principles, the majority of new corporate databases are still being created and managed with an RDBMS.

SQL Statements

1. Data Retrieval(DR)
2. Data Manipulation Language(DML)
3. Data Definition Language(DDL)
4. Data Control Language(DCL)
5. Transaction Control Language(TCL)

TYPE	STATEMENT	DESCRIPTION
DR	SELECT	Retrieves the data from the database
DML	1. INSERT 2. UPDATE 3. DELETE 4. MERGE	Enter new rows, changes existing rows, removes unwanted rows from tables in the database respectively.
DDL	1. CREATE 2. ALTER 3. DROP 4. RENAME 5. TRUNCATE	Sets up, changes and removes data structures from tables.
TCL	1. COMMIT 2. ROLLBACK 3. SAVEPOINT	Manages the changes made by DML statements. Changes to the data can be grouped together into logical transactions.
DCL	1. GRANT 2. RREVOKE	Gives or removes access rights to both the oracle database and the structures within it.

DATA TYPES

1. Character Data types:

- Char – fixed length character string that can varies between 1-2000 bytes
- Varchar / Varchar2 – variable length character string, size ranges from 1-4000 bytes.it saves the disk space(only length of the entered value will be assigned as the size of column)
- Long - variable length character string, maximum size is 2 GB

2. Number Data types : Can store +ve,-ve,zero,fixed point, floating point with 38 precession.

- Number – {p=38,s=0}
- Number(p) - fixed point
- Number(p,s) –floating point (p=1 to 38,s= -84 to 127)

3. Date Time Data type: used to store date and time in the table.

- DB uses its own format of storing in fixed length of 7 bytes for century, date, month, year, hour, minutes, and seconds.
- Default data type is “dd-mon-yy”
- New Date time data types have been introduced. They are
TIMESTAMP-Date with fractional seconds
INTERVAL YEAR TO MONTH-stored as an interval of years and months

INTERVAL DAY TO SECOND-stored as o interval of days to hour’s minutes and seconds

4. **Raw Data type:** used to store byte oriented data like binary data and byte string.

5. **Other :**

- CLOB – stores character object with single byte character.
- BLOB – stores large binary objects such as graphics, video, sounds.
- BFILE – stores file pointers to the LOB’s.

EXERCISE-1

Creating and Managing Tables

OBJECTIVE

After the completion of this exercise, students should be able to do the following:

- Create tables
- Describing the data types that can be used when specifying column definition
- Alter table definitions
- Drop, rename, and truncate tables

NAMING RULES

Table names and column names:

- Must begin with a letter
- Must be 1-30 characters long
- Must contain only A-Z, a-z, 0-9, _, \$, and #
- Must not duplicate the name of another object owned by the same user
- Must not be an oracle server reserve words
- 2 different tables should not have same name.
- Should specify a unique column name.
- Should specify proper data type along with width
- Can include “not null” condition when needed. By default it is ‘null’.

The CREATE TABLE Statement

Table: Basic unit of storage; composed of rows and columns

Syntax: 1 Create table table_name (column_name1 data_type (size)
column_name2 data_type (size)...);

Syntax: 2 Create table table_name (column_name1 data_type (size) constraints,
column_name2 data_type constraints ...);

Example:

Create table employees (employee_id number(6), first_name varchar2(20), ..job_id varchar2(10),
CONSTRAINT emp_emp_id_pk PRIMARY KEY (employee_id));

Tables Used in this course

Creating a table by using a Sub query

SYNTAX

// CREATE TABLE table_name(column_name type(size)...);

Create table table_name as select column_name1,column_name2,.....colmn_namen from

table_name where predicate;

AS Subquery

Subquery is the select statement that defines the set of rows to be inserted into the new table.

Example

Create table dept80 as select employee_id, last_name, salary*12 Annsal, hire_date
from employees where dept_id=80;

The ALTER TABLE Statement

The ALTER statement is used to

- Add a new column
- Modify an existing column
- Define a default value to the new column
- Drop a column
- To include or drop integrity constraint.

SYNTAX

ALTER TABLE table_name ADD /MODIFY(Column_name type(size));

ALTER TABLE table_name DROP COLUMN (Column_nname);

ALTER TABLE ADD CONSTRAINT Constraint_name PRIMARY KEY (Colum_Name);

Example:

Alter table dept80 add (jod_id varchar2(9));

Alter table dept80 modify (last_name varchar2(30));

Alter table dept80 drop column job_id;

NOTE: Once the column is dropped it cannot be recovered.

DROPPING A TABLE

- All data and structure in the table is deleted.
- Any pending transactions are committed.
- All indexes are dropped.
- Cannot roll back the drop table statement.

Syntax:

Drop table *tablename*;

Example:

Drop table dept80;

RENAMING A TABLE

To rename a table or view.

Syntax

RENAME old_name to new_name

Example:

Rename dept to detail_dept;

TRUNCATING A TABLE

Removes all rows from the table.

Releases the storage space used by that table.

Syntax

TRUNCATE TABLE *table_name*;

Example:

TRUNCATE TABLE copy_emp;

Find the Solution for the following:

Create the following tables with the given structure.

EMPLOYEES TABLE

NAME	NULL?	TYPE
Employee_id	Not null	Number(6)
First_Name		Varchar(20)
Last_Name	Not null	Varchar(25)
Email	Not null	Varchar(25)
Phone_Number		Varchar(20)
Hire_date	Not null	Date
Job_id	Not null	Varchar(10)
Salary		Number(8,2)
Commission_pct		Number(2,2)
Manager_id		Number(6)
Department_id		Number(4)

DEPARTMENT TABLE

NAME	NULL?	TYPE
Dept_id	Not null	Number(6)
Dept_name	Not null	Varchar(20)
Manager_id		Number(6)
Location_id		Number(4)

JOB_GRADE TABLE

NAME	NULL?	TYPE
------	-------	------

Grade_level		Varchar(2)
Lowest_sal		Number
Highest_sal		Number

LOCATION TABLE

NAME	NULL?	TYPE
Location_id	Not null	Number(4)
St_addr		Varchar(40)
Postal_code		Varchar(12)
City	Not null	Varchar(30)
State_province		Varchar(25)
Country_id		Char(2)

1. Create the DEPT table based on the DEPARTMENT following the table instance chart below. Confirm that the table is created.

Column name	ID	NAME
Key Type		
Nulls/Unique		
FK table		
FK column		
Data Type	Number	Varchar2
Length	7	25

```
CREATE TABLE DEPT (
  ID NUMBER(7),
  NAME VARCHAR2(25));
```

2. Create the EMP table based on the following instance chart. Confirm that the table is created.

Column name	ID	LAST_NAME	FIRST_NAME	DEPT_ID
Key Type				
Nulls/Unique				
FK table				
FK column				
Data Type	Number	Varchar2	Varchar2	Number
Length	7	25	25	7

```
CREATE TABLE EMP (
  ID NUMBER(7), LAST_NAME VARCHAR2(25), FIRST_NAME VARCHAR2(25),
  DEPT_ID NUMBER(7));
```

3. Modify the EMP table to allow for longer employee last names. Confirm the modification.(Hint: Increase the size to 50)

```
ALTER TABLE EMP MODIFY (LAST_NAME VARCHAR2(50));
```

4. Create the EMPLOYEES2 table based on the structure of EMPLOYEES table. Include Only the Employee_id, First_name, Last_name, Salary and Dept_id coloumns. Name the columns Id, First_name, Last_name, salary and Dept_id respectively.

```
CREATE TABLE EMPLOYEES2 (  
    ID NUMBER(7), FIRST_NAME VARCHAR2(25), LAST_NAME VARCHAR2(25), SALARY  
    NUMBER(9), DEPT_ID NUMBER(7));
```

5. Drop the EMP table.

```
DROP TABLE EMP;
```

6. Rename the EMPLOYEES2 table as EMP.

```
RENAME EMPLOYEES2 TO EMP;
```

7. Add a comment on DEPT and EMP tables. Confirm the modification by describing the table.

```
COMMENT ON TABLE DEPT is 'This is a table';  
DESC emp;
```

8. Drop the First_name column from the EMP table and confirm it.

```
ALTER TABLE emp  
DROP COLUMN first_name;
```

Evaluation Procedure	Marks awarded
Query(5)	
Execution (5)	
Viva(5)	
Total (15)	
Faculty Signature	

EXERCISE-2

MANIPULATING DATA

OBJECTIVE

After, the completion of this exercise the students will be able to do the following

- Describe each DML statement
- Insert rows into tables
- Update rows into table
- Delete rows from table
- Control Transactions

A DML statement is executed when you:

- Add new rows to a table
- Modify existing rows
- Removing existing rows

A transaction consists of a collection of DML statements that form a logical unit of work.

To Add a New Row

INSERT Statement

Syntax

INSERT INTO table_name VALUES (column1 values, column2 values, ..., columnn values);

Example:

INSERT INTO department (70, 'Public relations', 100,1700);

Inserting rows with null values

Implicit Method: (Omit the column)

INSERT INTO department VALUES (30,'purchasing');

Explicit Method: (Specify NULL keyword)

INSERT INTO department VALUES (100,'finance', NULL, NULL);

Inserting Special Values

Example:

Using SYSDATE

INSERT INTO employees VALUES (113,'louis', 'popp', 'lpopp', '5151244567',**SYSDATE**, 'ac_account', 6900, NULL, 205, 100);

Inserting Specific Date Values

Example:

```
INSERT INTO employees VALUES ( 114,'den', 'raphealy', 'drapheal', '5151274561',  
TO_DATE('feb 3,1999','mon, dd ,yyyy'), 'ac_account', 11000,100,30);
```

To Insert Multiple Rows

& is the placeholder for the variable value

Example:

```
INSERT INTO department VALUES (&dept_id, &dept_name, &location);
```

Copying Rows from another table

➤ Using Subquery

Example:

```
INSERT INTO sales_reps(id, name, salary, commission_pct)  
SELECT employee_id, Last_name, salary, commission_pct  
FROM employees  
WHERE job_id LIKE '%REP');
```

CHANGING DATA IN A TABLE

UPDATE Statement

Syntax1: (to update specific rows)

```
UPDATE table_name SET column=value WHERE condition;
```

Syntax 2: (To update all rows)

```
UPDATE table_name SET column=value;
```

Updating columns with a subquery

```
UPDATE employees  
SET job_id= (SELECT job_id  
FROM employees  
WHERE employee_id=205)  
WHERE employee_id=114;
```

REMOVING A ROW FROM A TABLE

DELETE STATEMENT

Syntax

```
DELETE FROM table_name WHERE conditions;
```

Example:

```
DELETE FROM department WHERE dept_name='finance';
```

Find the Solution for the following:

1. Create MY_EMPLOYEE table with the following structure

NAME	NULL?	TYPE
ID	Not null	Number(4)
Last_name		Varchar(25)
First_name		Varchar(25)
Userid		Varchar(25)
Salary		Number(9,2)

2. Add the first and second rows data to MY_EMPLOYEE table from the following sample data.

ID	Last_name	First_name	Userid	salary
1	Patel	Ralph	rpatel	895
2	Dancs	Betty	bdancs	860
3	Biri	Ben	bbiri	1100
4	Newman	Chad	Cnewman	750
5	Ropebur	Audrey	aropebur	1550

3. Display the table with values.

```
SELECT * FROM MY_EMPLOYEE;
```

4. Populate the next two rows of data from the sample data. Concatenate the first letter of the first_name with the first seven characters of the last_name to produce Userid.

```
INSERT INTO MY_EMPLOYEE (ID, Last_name, First_name, Userid, Salary) VALUES (3, 'Biri', 'Ben', 'bbiri', 1100);  
INSERT INTO MY_EMPLOYEE (ID, Last_name, First_name, Userid, Salary) VALUES (4, 'Newman', 'Chad',  
'Cnewman', 750);
```

5. Make the data additions permanent.

```
COMMIT;
```

6. Change the last name of employee 3 to Drexler.

```
UPDATE MY_EMPLOYEE  
SET Last_name = 'Drexler'  
WHERE ID = 3;
```

7. Change the salary to 1000 for all the employees with a salary less than 900.

```
UPDATE MY_EMPLOYEE  
SET Salary = 1000  
WHERE Salary < 900;
```

8. Delete Betty dancs from MY _EMPLOYEE table.

```
DELETE FROM MY_EMPLOYEE  
WHERE First_name = 'Betty' AND Last_name = 'Dancs';
```

9. Empty the fourth row of the emp table.

```
DELETE FROM MY_EMPLOYEE  
WHERE ID = 4;
```

Evaluation Procedure	Marks awarded
Query(5)	
Execution (5)	
Viva(5)	
Total (15)	
Faculty Signature	

EXERCISE-3

INCLUDING CONSTRAINTS

OBJECTIVE

After the completion of this exercise the students should be able to do the following

- Describe the constraints
- Create and maintain the constraints

What are Integrity constraints?

- Constraints enforce rules at the table level.
- Constraints prevent the deletion of a table if there are dependencies

The following types of integrity constraints are valid

a) **Domain Integrity**

- ✓ NOT NULL
- ✓ CHECK

b) **Entity Integrity**

- ✓ UNIQUE
- ✓ PRIMARY KEY

c) **Referential Integrity**

- ✓ FOREIGN KEY

Constraints can be created in either of two ways

1. At the same time as the table is created
2. After the table has been created.

Defining Constraints

Create table tablename (column_name1 data_type constraints, column_name2 data_type constraints ...);

Example:

Create table employees (employee_id number(6), first_name varchar2(20), ..job_id varchar2 (10),
CONSTRAINT emp_emp_id_pk PRIMARY KEY (employee_id));

Domain Integrity

This constraint sets a range and any violations that takes place will prevent the user from performing the manipulation that caused the breach.It includes:

NOT NULL Constraint

While creating tables, by default the rows can have null value.the enforcement of not null constraint in a table ensure that the table contains values.

Principle of null values:

- Setting null value is appropriate when the actual value is unknown, or when a value would not be meaningful.
- A null value is not equivalent to a value of zero.
- A null value will always evaluate to null in any expression.
- When a column name is defined as not null, that column becomes a mandatory i.e., the user has to enter data into it.
- Not null Integrity constraint cannot be defined using the alter table command when the table contain rows.

Example

```
CREATE TABLE employees (employee_id number (6), last_name varchar2(25) NOT NULL, salary number(8,2), commission_pct number(2,2), hire_date date constraint emp_hire_date_nn NOT NULL'....);
```

CHECK

Check constraint can be defined to allow only a particular range of values.when the manipulation violates this constraint,the record will be rejected.Check condition cannot contain sub queries.

```
CREATE TABLE employees (employee_id number (6), last_name varchar2 (25) NOT NULL, salary number(8,2), commission_pct number(2,2), hire_date date constraint emp_hire_date_nn NOT NULL'...,CONSTRAINT emp_salary_mi CHECK(salary > 0));
```

Entity Integrity

Maintains uniqueness in a record. An entity represents a table and each row of a table represents an instance of that entity. To identify each row in a table uniquely we need to use this constraint.

There are 2 entity constraints:

a) Unique key constraint

It is used to ensure that information in the column for each record is unique, as with telephone or driver's license numbers. It prevents the duplication of value with rows of a specified column in a set of column. A column defined with the constraint can allow null value.

If unique key constraint is defined in more than one column i.e., combination of column cannot be specified. Maximum combination of columns that a composite unique key can contain is 16.

Example:

```
CREATE TABLE employees (employee_id number(6), last_name varchar2(25) NOT NULL,email varchar2(25), salary number(8,2), commission_pct number(2,2), hire_date date constraint emp_hire_date_nn NOT NULL' CONSTRAINT emp_email_uk UNIQUE(email));
```

PRIMARY KEY CONSTRAINT

A primary key avoids duplication of rows and does not allow null values. Can be defined on one or more columns in a table and is used to uniquely identify each row in a table. These values should never be changed and should never be null.

A table should have only one primary key. If a primary key constraint is assigned to more than one column or combination of column is said to be composite primary key, which can contain 16 columns.

Example:

```
CREATE TABLE employees (employee_id number(6) , last_name varchar2(25) NOT NULL, email varchar2(25), salary number(8,2), commission_pct number(2,2), hire_date date constraint emp_hire_date_nn NOT NULL, Constraint emp_id pk PRIMARY KEY (employee_id), CONSTRAINT emp_email_uk UNIQUE(email));
```

c) Referential Integrity

It enforces relationship between tables. To establish parent-child relationship between 2 tables having a common column definition, we make use of this constraint. To implement this, we should define the column in the parent table as primary key and same column in the child table as foreign key referring to the corresponding parent entry.

Foreign key

A column or combination of column included in the definition of referential integrity, which would refer to a referenced key.

Referenced key

It is a unique or primary key upon which is defined on a column belonging to the parent table.

Keywords:

FOREIGN KEY: Defines the column in the child table at the table level constraint.

REFERENCES: Identifies the table and column in the parent table.

ON DELETE CASCADE: Deletes the dependent rows in the child table when a row in the parent table is deleted.

ON DELETE SET NULL: converts dependent foreign key values to null when the parent value is removed.

```
CREATE TABLE employees (employee_id number(6) , last_name varchar2(25) NOT NULL, email varchar2(25), salary number(8,2), commission_pct number(2,2), hire_date date constraint emp_hire_date_nn NOT NULL, Constraint emp_id pk PRIMARY KEY (employee_id), CONSTRAINT emp_email_uk UNIQUE(email), CONSTRAINT emp_dept_fk FOREIGN KEY (department_id) references departments(dept_id));
```

ADDING A CONSTRAINT

Use the ALTER to

- Add or Drop a constraint, but not modify the structure
- Enable or Disable the constraints
- Add a not null constraint by using the Modify clause

Syntax

ALTER TABLE table name ADD CONSTRAINT Cons_name type(column name);

Example:

ALTER TABLE employees ADD CONSTRAINT emp_manager_fk FOREIGN KEY (manager_id) REFERENCES employees (employee_id);

DROPPING A CONSTRAINT

Example:

ALTER TABLE employees DROP CONSTRAINT emp_manager_fk;

CASCADE IN DROP

- The CASCADE option of the DROP clause causes any dependent constraints also to be dropped.

Syntax

ALTER TABLE departments DROP PRIMARY KEY|UNIQUE (column)| CONSTRAINT constraint _name CASCADE;

DISABLING CONSTRAINTS

- Execute the DISABLE clause of the ALTER TABLE statement to deactivate an integrity constraint
- Apply the CASCADE option to disable dependent integrity constraints.

Example

ALTER TABLE employees DISABLE CONSTRAINT emp_emp_id_pk CASCADE;

ENABLING CONSTRAINTS

- Activate an integrity constraint currently disabled in the table definition by using the ENABLE clause.

Example

ALTER TABLE employees ENABLE CONSTRAINT emp_emp_id_pk CASCADE;

CASCADING CONSTRAINTS

The CASCADE CONSTRAINTS clause is used along with the DROP column clause. It drops all referential integrity constraints that refer to the primary and unique keys defined on the dropped Columns.

This clause also drops all multicolumn constraints defined on the dropped column.

Example:

Assume table TEST1 with the following structure

```
CREATE TABLE test1 ( pk number PRIMARY KEY, fk number, col1 number,col2 number,  
CONSTRAINT fk_constraint FOREIGN KEY(fk) references test1, CONSTRAINT ck1 CHECK  
(pk>0 and col1>0), CONSTRAINT ck2 CHECK (col2>0));
```

An error is returned for the following statements

```
ALTER TABLE test1 DROP (pk);
```

```
ALTER TABLE test1 DROP (col1);
```

The above statement can be written with CASCADE CONSTRAINT

```
ALTER TABLE test 1 DROP(pk) CASCADE CONSTRAINTS;
```

(OR)

```
ALTER TABLE test 1 DROP(pk, fk, col1) CASCADE CONSTRAINTS;
```

VIEWING CONSTRAINTS

Query the USER_CONSTRAINTS table to view all the constraints definition and names.

Example:

```
SELECT constraint_name, constraint_type, search_condition FROM user_constraints  
WHERE table_name='employees';
```

Viewing the columns associated with constraints

```
SELECT constraint_name, constraint_type, FROM user_cons_columns  
WHERE table_name='employees';
```

Find the Solution for the following:

1. Add a table-level PRIMARY KEY constraint to the EMP table on the ID column. The constraint should be named at creation. Name the constraint my_emp_id_pk.

```
ALTER TABLE EMP  
ADD CONSTRAINT my_emp_id_pk PRIMARY KEY (ID);
```

2. Create a PRIMAY KEY constraint to the DEPT table using the ID colum. The constraint should be named at creation. Name the constraint my_dept_id_pk.

```
ALTER TABLE DEPT  
ADD CONSTRAINT my_dept_id_pk PRIMARY KEY (ID);
```

3. Add a column DEPT_ID to the EMP table. Add a foreign key reference on the EMP table that ensures that the employee is not assigned to nonexistent department. Name the constraint my_emp_dept_id_fk.

```
ALTER TABLE EMP  
ADD DEPT_ID NUMBER;
```

```
ALTER TABLE EMP  
ADD CONSTRAINT my_emp_dept_id_fk FOREIGN KEY (DEPT_ID)  
REFERENCES DEPT(ID);
```

4. Modify the EMP table. Add a COMMISSION column of NUMBER data type, precision 2, scale 2. Add a constraint to the commission column that ensures that a commission value is greater than zero.

```
ALTER TABLE EMP  
ADD COMMISSION NUMBER(5, 2);  
ALTER TABLE EMP  
ADD CONSTRAINT chk_commission_positive CHECK (COMMISSION > 0);
```

Evaluation Procedure	Marks awarded
Query(5)	
Execution (5)	
Viva(5)	
Total (15)	
Faculty Signature	

EXERCISE-4

Writing Basic SQL SELECT Statements

OBJECTIVES

After the completion of this exercise, the students will be able to do the following:

- List the capabilities of SQL SELECT Statement
- Execute a basic SELECT statement

Capabilities of SQL SELECT statement

A SELECT statement retrieves information from the database. Using a select statement, we can perform

- ✓ Projection: To choose the columns in a table
- ✓ Selection: To choose the rows in a table
- ✓ Joining: To bring together the data that is stored in different tables

Basic SELECT Statement

Syntax

```
SELECT *|DISTINCT Column_ name| alias  
FROM table_name;
```

NOTE:

DISTINCT—Suppr
ess the duplicates.

Alias—gives selected columns different headings.

Example: 1

```
SELECT * FROM departments;
```

Example: 2

```
SELECT location_id, department_id FROM departments;
```

Writing SQL Statements

- SQL statements are not case sensitive
- SQL statements can be on one or more lines.
- Keywords cannot be abbreviated or split across lines
- Clauses are usually placed on separate lines
- Indents are sued to enhance readability

Using Arithmetic Expressions

Basic Arithmetic operators like *, /, +, -can be used

Example:1

```
SELECT last_name, salary, salary+300 FROM employees;
```

Example:2

```
SELECT last_name, salary, 12*salary+100 FROM employees;
```

The statement is not same as

```
SELECT last_name, salary, 12*(salary+100) FROM employees;
```

Example:3

SELECT last_name, job_id, salary, commission_pct FROM employees;

Example:4

SELECT last_name, job_id, salary, 12*salary*commission_pct FROM employees;

Using Column Alias

- To rename a column heading with or without AS keyword.

Example:1

SELECT last_name AS Name
FROM employees;

Example: 2

SELECT last_name "Name" salary*12 "Annual Salary "
FROM employees;

Concatenation Operator

- Concatenates columns or character strings to other columns
- Represented by two vertical bars (||)
- Creates a resultant column that is a character expression

Example:

SELECT last_name||job_id AS "EMPLOYEES JOB" FROM employees;

Using Literal Character String

- A literal is a character, a number, or a date included in the SELECT list.
- Date and character literal values must be enclosed within single quotation marks.

Example:

SELECT last_name||'is a'||job_id AS "EMPLOYEES JOB" FROM employees;

Eliminating Duplicate Rows

- Using DISTINCT keyword.

Example:

SELECT DISTINCT department_id FROM employees;

Displaying Table Structure

- Using DESC keyword.

Syntax

DESC table_name;

Example:

DESC employees;

Find the Solution for the following:

True OR False

1. The following statement executes successfully.

Identify the Errors

```
SELECT employee_id, last_name  
sal*12 ANNUAL SALARY  
FROM employees;
```

Queries

```
SELECT employee_id, last_name, sal * 12 AS "ANNUAL SALARY"  
FROM employees;
```

2. Show the structure of departments the table. Select all the data from it.

```
DESC DEPARTMENTS;
```

3. Create a query to display the last name, job code, hire date, and employee number for each employee, with employee number appearing first.

```
SELECT employee_id AS "EMPLOYEE NUMBER", last_name, job_id AS "JOB CODE",  
hire_date AS STARTDATE FROM employees;
```

4. Provide an alias STARTDATE for the hire date.

```
SELECT employee_id AS "EMPLOYEE NUMBER", last_name, job_id AS "JOB CODE",  
hire_date AS STARTDATE  
FROM employees;
```

5. Create a query to display unique job codes from the employee table.

```
SELECT DISTINCT job_id  
FROM employees;
```

6. Display the last name concatenated with the job ID , separated by a comma and space, and name the column EMPLOYEE and TITLE.

```
SELECT last_name || ', ' || job_id AS "EMPLOYEE and TITLE"  
FROM employees;
```

7. Create a query to display all the data from the employees table. Separate each column by a comma. Name the column THE_OUTPUT.

```
SELECT *  
FROM employees;
```

Evaluation Procedure	Marks awarded
Query(5)	
Execution (5)	
Viva(5)	
Total (15)	
Faculty Signature	

EXERCISE-5

Restricting and Sorting data

After the completion of this exercise, the students will be able to do the following:

- Limit the rows retrieved by the queries
- Sort the rows retrieved by the queries
-

Limiting the Rows selected

- Using WHERE clause
- Alias cannot be used in WHERE clause

Syntax

SELECT-----
FROM-----
WHERE condition;

Example:

SELECT employee_id, last_name, job_id, department_id FROM employees WHERE
department_id=90;

Character strings and Dates

Character strings and date values are enclosed in single quotation marks.

Character values are case sensitive and date values are format sensitive.

Example:

SELECT employee_id, last_name, job_id, department_id FROM employees
WHERE last_name='WHALEN';

Comparison Conditions

All relational operators can be used. (=, >, >=, <, <=, <>, !=)

Example:

SELECT last_name, salary
FROM employees
WHERE salary<=3000;

Other comparison conditions

Operator	Meaning
BETWEEN ...AND...	Between two values
IN	Match any of a list of values
LIKE	Match a character pattern
IS NULL	Is a null value

Example:1

```
SELECT last_name, salary
FROM employees
WHERE salary BETWEEN 2500 AND 3500;
```

Example:2

```
SELECT employee_id, last_name, salary , manager_id
FROM employees
WHERE manager_id IN (101, 100,201);
```

Example:3

- Use the LIKE condition to perform wildcard searches of valid string values.
- Two symbols can be used to construct the search string
 - % denotes zero or more characters
 - _ denotes one character

```
SELECT first_name, salary
FROM employees
WHERE first_name LIKE '%s';
```

Example:4

```
SELECT last_name, salary
FROM employees
WHERE last_name LIKE '_o%';
```

Example:5

ESCAPE option-To have an exact match for the actual % and _ characters
To search for the string that contain 'SA_'

```
SELECT employee_id, first_name, salary, job_id
FROM employees
WHERE job_id LIKE '%sa\_%' ESCAPE '\';
```

Test for NULL

- Using IS NULL operator

Example:

```
SELECT employee_id, last_name, salary , manager_id
FROM employees
WHERE manager_id IS NULL;
```

Logical Conditions

All logical operators can be used.(AND,OR,NOT)

Example:1

```
SELECT employee_id, last_name, salary , job_id
```

```
FROM employees
WHERE salary>=10000
AND job_id LIKE '%MAN%';
```

Example:2

```
SELECT employee_id, last_name, salary , job_id
FROM employees
WHERE salary>=10000
OR job_id LIKE '%MAN%';
```

Example:3

```
SELECT employee_id, last_name, salary , job_id
FROM employees
WHERE job_id NOT IN ('it_prog', st_clerk', sa_rep');
```

Rules of Precedence

Order Evaluated	Operator
1	Arithmetic
2	Concatenation
3	Comparison
4	IS [NOT] NULL, LIKE, [NOT] IN
5	[NOT] BETWEEN
6	Logical NOT
7	Logical AND
8	Logical OR

Example:1

```
SELECT employee_id, last_name, salary , job_id
FROM employees
WHERE job_id ='sa_rep'
OR job_id='ad_pres'
AND salary>15000;
```

Example:2

```
SELECT employee_id, last_name, salary , job_id
FROM employees
WHERE (job_id ='sa_rep'
OR job_id='ad_pres')
AND salary>15000;
```

Sorting the rows

Using ORDER BY Clause

ASC-Ascending Order,Default

DESC-Descending order

Example:1

```
SELECT last_name, salary , job_id,department_id,hire_date
FROM employees
ORDER BY hire_date;
```

Example:2

```
SELECT last_name, salary , job_id,department_id,hire_date
FROM employees
ORDER BY hire_date DESC;
```

Example:3

Sorting by column alias

```
SELECT last_name, salary*12 annsal , job_id,department_id,hire_date
FROM employees
ORDER BY annsal;
```

Example:4

Sorting by Multiple columns

```
SELECT last_name, salary , job_id,department_id,hire_date
FROM employees
ORDER BY department_id, salary DESC;
```

Find the Solution for the following:

1. Create a query to display the last name and salary of employees earning more than 12000.

```
SELECT last_name, salary
FROM employees
WHERE salary > 12000;
```

2. Create a query to display the employee last name and department number for employee number 176.

```
SELECT last_name, department_id
FROM employees
WHERE employee_id = 176;
```

3. Create a query to display the last name and salary of employees whose salary is not in the range of 5000 and 12000. (hints: not between)

```
SELECT last_name, salary
FROM employees
WHERE salary NOT BETWEEN 5000 AND 12000;
```

4. Display the employee last name, job ID, and start date of employees hired between February 20,1998 and May 1,1998.order the query in ascending order by start date.(hints: between)

```
SELECT last_name, job_id, hire_date
FROM employees
WHERE hire_date BETWEEN TO_DATE('1998-02-20', 'YYYY-MM-DD') AND
TO_DATE('1998-05-01', 'YYYY-MM-DD')
ORDER BY hire_date ASC;
```

5. Display the last name and department number of all employees in departments 20 and 50 in alphabetical order by name.(hints: in, orderby)

```
SELECT last_name, department_id
FROM employees
WHERE department_id IN (20, 50)
ORDER BY last_name;
```

6. Display the last name and salary of all employees who earn between 5000 and 12000 and are in departments 20 and 50 in alphabetical order by name. Label the columns EMPLOYEE, MONTHLY SALARY respectively.(hints: between, in)

```
SELECT last_name AS EMPLOYEE, salary AS "MONTHLY SALARY"
FROM employees
WHERE department_id IN (20, 50)
AND salary BETWEEN 5000 AND 12000
ORDER BY last_name;
```

7. Display the last name and hire date of every employee who was hired in 1994.(hints: like)

```
SELECT last_name, hire_date
FROM employees
```

```
WHERE hire_date LIKE '1994%';
```

8. Display the last name and job title of all employees who do not have a manager.(hints: is null)

```
SELECT last_name, job_title  
FROM employees  
WHERE manager_id IS NULL;
```

9. Display the last name, salary, and commission for all employees who earn commissions. Sort data in descending order of salary and commissions.(hints: is not nul,orderby)

```
SELECT last_name, salary, commission_pct AS commission  
FROM employees  
WHERE commission_pct IS NOT NULL  
ORDER BY salary DESC, commission_pct DESC;
```

10. Display the last name of all employees where the third letter of the name is *a*.(hints:like)

```
SELECT last_name  
FROM employees  
WHERE SUBSTR(last_name, 3, 1) = 'a';
```

11. Display the last name of all employees who have an *a* and an *e* in their last name.(hints: like)

```
SELECT last_name  
FROM employees  
WHERE last_name LIKE '%a%' AND last_name LIKE '%e%';
```

12. Display the last name and job and salary for all employees whose job is sales representative or stock clerk and whose salary is not equal to 2500 ,3500 or 7000.(hints:in,not in)

```
SELECT last_name, job_id, salary  
FROM employees  
WHERE job_id IN ('SA_REP', 'ST_CLERK')  
AND salary NOT IN (2500, 3500, 7000);
```

13. Display the last name, salary, and commission for all employees whose commission amount is 20%.(hints:use predicate logic)

```
SELECT last_name, salary, commission_pct AS commission
FROM employees
WHERE commission_pct = 0.20;
```

Evaluation Procedure	Marks awarded
Query(5)	
Execution (5)	
Viva(5)	
Total (15)	
Faculty Signature	

EXERCISE-6

Single Row Functions

Objective

After the completion of this exercise, the students will be able to do the following:

- Describe various types of functions available in SQL.
- Use character, number and date functions in SELECT statement.
- Describe the use of conversion functions.

Single row functions:

Manipulate data items.

Accept arguments and return one value.

Act on each row returned.

Return one result per row.

May modify the data type.

Can be nested.

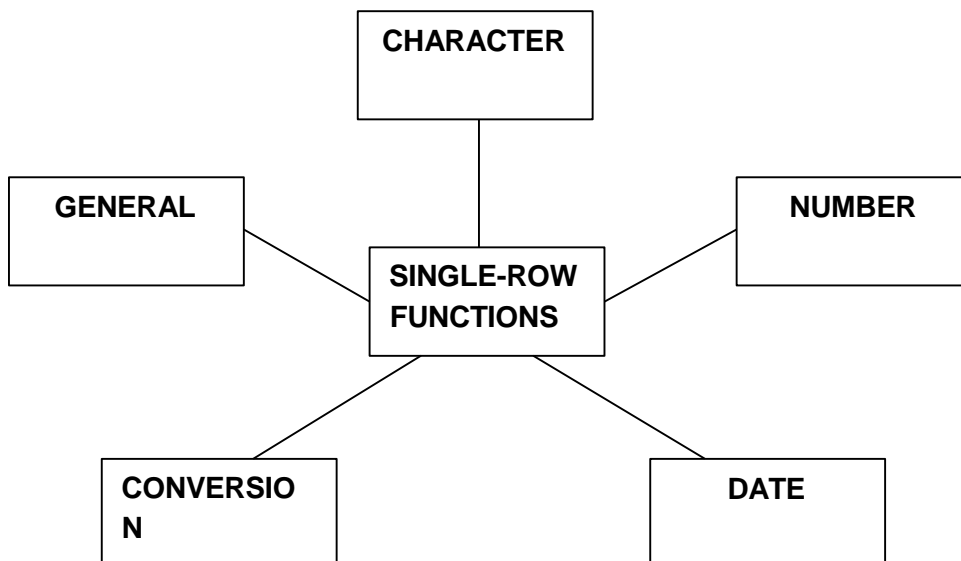
Accept arguments which can be a column or an expression

Syntax

Function_name(arg1,...argn)

An argument can be one of the following

- ✓ User-supplied constant
- ✓ Variable value
- ✓ Column name
- ✓ Expression



- Character Functions: Accept character input and can return both character and number values.
- Number functions: Accept numeric input and return numeric values.
- Date Functions: Operate on values of the DATE data type.
- Conversion Functions: Convert a value from one type to another.

Character Functions

Character Functions

Case-manipulation functions

1. Lower
2. Upper
3. Initcap

Character-manipulation functions

1. Concat
2. Substr
3. Length
4. Instr
5. Lpad/Rpad
6. Trim
7. Repalce

Function	Purpose
lower(column/expr)	Converts alpha character values to lowercase
upper(column/expr)	Converts alpha character values to uppercase
initcap(column/expr)	Converts alpha character values the to uppercase for the first letter of each word, all other letters in lowercase
concat(column1/expr1, column2/expr2)	Concatenates the first character to the second character
substr(column/expr,m,n)	Returns specified characters from character value starting at character position m, n characters long
length(column/expr)	Returns the number of characters in the expression
instr(column/expr,'string',m,n)	Returns the numeric position of a named string
lpad(column/expr, n,'string')	Pads the character value right-justified to a total width of n character positions
rpad(column/expr,'string',m,n)	Pads the character value left-justified to a total width of n character positions
trim(leading/trailing/both, trim_character FROM trim_source)	Enables you to trim heading or string. trailing or both from a character
replace(text, search_string, replacement_string)	

Example:

lower('SQL Course') sql course

□

upper('SQL Course') SQL COURSE

□

initcap('SQL Course') Sql Course

□

```
SELECT 'The job id for'|| upper(last_name||'is')||lower(job_id) AS "EMPLOYEE DETAILS"
FROM employees;
```

```
SELECT employee_id, last_name, department_id
FROM employees
WHERE LOWER(last_name)='higgins';
```

Function	Result
CONCAT('hello', 'world')	helloworld
Substr('helloworld',1,5)	Hello
Length('helloworld')	10
Instr('helloworld','w')	6
Lpad(salary,10,'*')	*****24000

Rpad(salary,10,'*')	24000*****
Trim('h' FROM 'helloworld')	elloworld

Command	Query	Output
initcap(char);	<i>select initcap("hello") from dual;</i>	Hello
lower (char); upper (char);	<i>select lower ('HELLO') from dual;</i> <i>select upper ('hello') from dual;</i>	Hello HELLO
ltrim (char,[set]);	<i>select ltrim ('cseit', 'cse') from dual;</i>	IT
rtrim (char,[set]);	<i>select rtrim ('cseit', 'it') from dual;</i>	CSE
replace (char,search string, replace string);	<i>select replace ('jack and jue', 'j', 'bl') from dual;</i>	black and blue
substr (char,m,n);	<i>select substr ('information', 3, 4) from dual;</i>	form

Example:

SELECT employee_id, CONCAT (first_name,last_name) NAME , job_id,LENGTH(last_name),
INSTR(last_name,'a') "contains'a'?"
FROM employees WHERE SUBSTR(job_id,4)='ERP';

NUMBER FUNCTIONS

Function	Purpose
round(column/expr, n)	Rounds the value to specified decimal
trunc(column/expr,n)	Truncates value to specified decimal
mod(m,n)	Returns remainder of division

Example

Function	Result
round(45.926,2)	45.93
trunc(45.926,2)	45.92
mod(1600,300)	100

SELECT ROUND(45.923,2), ROUND(45.923,0), ROUND(45.923,-1) FROM dual;

NOTE: Dual is a dummy table you can use to view results from functions and calculations.

SELECT TRUNC(45.923,2), TRUNC(45.923), TRUNC(45.923,-2) FROM dual;

SELECT last_name,salary,MOD(salary,5000) FROM employees WHERE job_id='sa_rep';

Working with Dates

The Oracle database stores dates in an internal numeric format: century, year, month, day, hours, minutes, and seconds.

- The default date display format is DD-MON-RR.
- Enables you to store 21st-century dates in the 20th century by specifying only the last two digits of the year
- Enables you to store 20th-century dates in the 21st century in the same way

Example

SELECT last_name, hire_date FROM employees WHERE hire_date < '01-FEB-88;

Working with Dates

SYSDATE is a function that returns:

- Date
- Time

Example

Display the current date using the DUAL table.

```
SELECT SYSDATE FROM DUAL;
```

Arithmetic with Dates

- Add or subtract a number to or from a date for a resultant date value.
- Subtract two dates to find the number of days between those dates.
- Add hours to a date by dividing the number of hours by 24.

Arithmetic with Dates

Because the database stores dates as numbers, you can perform calculations using arithmetic Operators such as addition and subtraction. You can add and subtract number constants as well as dates.

You can perform the following operations:

Operation	Result	Description
date + number	Date	Adds a number of days to a date
date – number	Date	Subtracts a number of days from a date
date – date	Number of days	Subtracts one date from another
date + number/24	Date	Adds a number of hours to a date

Example

```
SELECT last_name, (SYSDATE-hire_date)/7 AS WEEKS  
FROM employees  
WHERE department_id = 90;
```

Date Functions

Function	Result
MONTHS_BETWEEN	Number of months between two dates
ADD_MONTHS	Add calendar months to date
NEXT_DAY	Next day of the date specified
LAST_DAY	Last day of the month
ROUND	Round date
TRUNC	Truncate date

Date Functions

Date functions operate on Oracle dates. All date functions return a value of DATE data type except MONTHS_BETWEEN, which returns a numeric value.

- MONTHS_BETWEEN(date1, date2):: Finds the number of months between date1 and date2. The result can be positive or negative. If date1 is later than date2, the result is positive; if date1 is earlier than date2, the result is negative. The noninteger part of the result represents a portion of the month.

- **ADD_MONTHS(date, n):::** Adds n number of calendar months to date. The value of n must be an integer and can be negative.
- **NEXT_DAY(date, 'char'):::** Finds the date of the next specified day of the week ('char') following date. The value of char may be a number representing a day or a character string.
- **LAST_DAY(date):::** Finds the date of the last day of the month that contains date
- **ROUND(date[, 'fmt']:::** Returns date rounded to the unit that is specified by the format model fmt. If the format model fmt is omitted, date is rounded to the nearest day.
- **TRUNC(date[, 'fmt']:::** Returns date with the time portion of the day truncated to the unit that is specified by the format model fmt. If the format model fmt is omitted, date is truncated to the nearest day.

Using Date Functions

Function	Result
MONTHS_BETWEEN ('01-SEP-95' , '11-JAN-94')	19.6774194
ADD_MONTHS ('11-JAN-94' , 6)	'11-JUL-94'
NEXT_DAY ('01-SEP-95' , 'FRIDAY')	'08-SEP-95'
LAST_DAY ('01-FEB-95')	'28-FEB-95'

Example


Display the employee number, hire date, number of months employed, sixmonth review date, first Friday after hire date, and last day of the hire month for all employees who have been employed for fewer than 70 months.

```
SELECT employee_id, hire_date, MONTHS_BETWEEN (SYSDATE, hire_date)
TENURE, ADD_MONTHS (hire_date, 6) REVIEW, NEXT_DAY (hire_date, 'FRIDAY'),
LAST_DAY(hire_date)
FROM employees
WHERE MONTHS_BETWEEN (SYSDATE, hire_date) < 70;
```

Conversion Functions

This covers the following topics:

- Writing a query that displays the current date
- Creating queries that require the use of numeric, character, and date functions
- Performing calculations of years and months of service for an employee



Implicit Data Type Conversion

For assignments, the Oracle server can automatically convert the following:

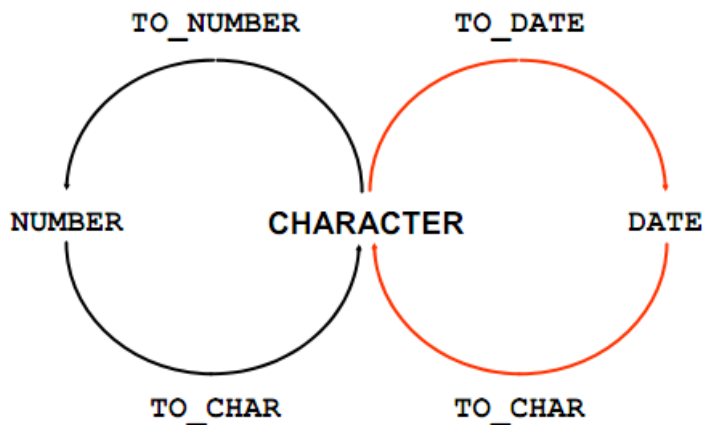
From	To
VARCHAR2 or CHAR	NUMBER
VARCHAR2 or CHAR	DATE
NUMBER	VARCHAR2
DATE	VARCHAR2

For example, the expression `hire_date > '01-JAN-90'` results in the implicit conversion from the string '01-JAN-90' to a date.

For expression evaluation, the Oracle Server can automatically convert the following:

From	To
VARCHAR2 or CHAR	NUMBER
VARCHAR2 or CHAR	DATE

Explicit Data Type Conversion



SQL provides three functions to convert a value from one data type to another:

Example:

Using the TO_CHAR Function with Dates

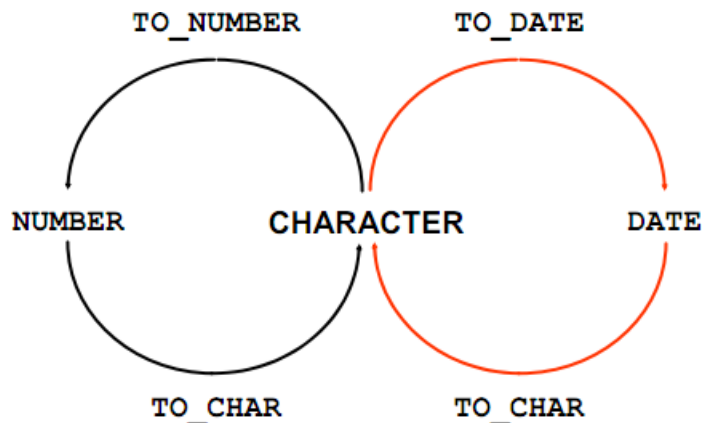
`TO_CHAR(date, 'format_model')`

The format model:

- Must be enclosed by single quotation marks
- Is case-sensitive
- Can include any valid date format element
- Has an fm element to remove padded blanks or suppress leading zeros
- Is separated from the date value by a comma

```
SELECT employee_id, TO_CHAR(hire_date, 'MM/YY') Month_Hired
FROM employees WHERE last_name = 'Higgins';
```

Elements of the Date Format Model



Sample Format Elements of Valid Date

Element	Description
SCC or CC	Century; server prefixes B.C. date with -
Years in dates YYYY or SYYYY	Year; server prefixes B.C. date with -
YYY or YY or Y	Last three, two, or one digits of year
Y,YYY	Year with comma in this position
IYYY, IYY, IY, I	Four-, three-, two-, or one-digit year based on the ISO standard
SYEAR or YEAR	Year spelled out; server prefixes B.C. date with -
BC or AD	Indicates B.C. or A.D. year
B.C. or A.D.	Indicates B.C. or A.D. year using periods
Q	Quarter of year
MM	Month: two-digit value
MONTH	Name of month padded with blanks to length of nine characters
MON	Name of month, three-letter abbreviation
RM	Roman numeral month
WW or W	Week of year or month
DDD or DD or D	Day of year, month, or week
DAY	Name of day padded with blanks to a length of nine characters
DY	Name of day; three-letter abbreviation
J	Julian day; the number of days since December 31, 4713 B.C.

Date Format Elements: Time Formats

Use the formats that are listed in the following tables to display time information and literals and to change numerals to spelled numbers.

Element	Description
AM or PM	Meridian indicator
A.M. or P.M.	Meridian indicator with periods
HH or HH12 or HH24	Hour of day, or hour (1–12), or hour (0–23)
MI	Minute (0–59)
SS	Second (0–59)
SSSSS	Seconds past midnight (0–86399)

Other Formats

Element	Description
/ . ,	Punctuation is reproduced in the result.
"of the"	Quoted string is reproduced in the result.

Specifying Suffixes to Influence Number Display

Element	Description
TH	Ordinal number (for example, DDTH for 4TH)
SP	Spelled-out number (for example, DDSP for FOUR)
SPTH or THSP	Spelled-out ordinal numbers (for example, DDSPTH for FOURTH)

Example

```
SELECT last_name,
       TO_CHAR(hire_date, 'fmDD Month YYYY') AS HIREDATE
FROM   employees;
```

Modify example to display the dates in a format that appears as “Seventeenth of June 1987 12:00:00 AM.”

```
SELECT last_name,
       TO_CHAR(hire_date, 'fmDdspt "of" Month YYYY fmHH:MI:SS AM') HIREDATE
FROM   employees;
```

Using the TO_CHAR Function with Numbers

TO_CHAR(number, 'format_model')

These are some of the format elements that you can use with the TO_CHAR function to display a number value as a character:

Element	Result
9	Represents a number
0	Forces a zero to be displayed
\$	Places a floating dollar sign
L	Uses the floating local currency symbol
.	Prints a decimal point
,	Prints a comma as thousands indicator

Number Format Elements

If you are converting a number to the character data type, you can use the following format elements:

Element	Description	Example	Result
9	Numeric position (number of 9s determine display width)	999999	1234
0	Display leading zeros	099999	001234
\$	Floating dollar sign	\$999999	\$1234
L	Floating local currency symbol	L999999	FF1234
D	Returns in the specified position the decimal character. The default is a period (.).	99D99	99.99
.	Decimal point in position specified	999999.99	1234.00
G	Returns the group separator in the specified position. You can specify multiple group separators in a number format model.	9,999	9G999
,	Comma in position specified	999,999	1,234
MI	Minus signs to right (negative values)	999999MI	1234-
PR	Parenthesize negative numbers	999999PR	<1234>
EEEE	Scientific notation (format must specify four Es)	99.999EEEE	1.234E+03
U	Returns in the specified position the "Euro" (or other) dual currency	U9999	€1234
V	Multiply by 10 <i>n</i> times (<i>n</i> = number of 9s after V)	9999V99	123400
S	Returns the negative or positive value	S9999	-1234 or +1234
B	Display zero values as blank, not 0	B9999.99	1234.00

```
SELECT TO_CHAR(salary, '$99,999.00') SALARY
FROM   employees
WHERE  last_name = 'Ernst';
```

Using the TO NUMBER and TO DATE Functions

- Convert a character string to a number format using the TO_NUMBER function:
TO_NUMBER(char[, 'format_model']
- Convert a character string to a date format using the TO_DATE function:
TO_DATE(char[, 'format_model']
- These functions have an fx modifier. This modifier specifies the exact matching for the character argument and date format model of a TO_DATE function.

The fx modifier specifies exact matching for the character argument and date format model of a TO_DATE function:

- Punctuation and quoted text in the character argument must exactly match (except for case) the corresponding parts of the format model.
- The character argument cannot have extra blanks. Without fx, Oracle ignores extra blanks.

- Numeric data in the character argument must have the same number of digits as the corresponding element in the format model. Without fx, numbers in the character argument can omit leading zeros.

```
SELECT last_name, hire_date
FROM employees
WHERE hire_date = TO_DATE('May 24, 1999', 'fxMonth DD, YYYY');
```

Find the Solution for the following:

1. Write a query to display the current date. Label the column Date.

```
SELECT CURRENT_DATE AS "Date";
```

2. The HR department needs a report to display the employee number, last name, salary, and increased by 15.5% (expressed as a whole number) for each employee. Label the column New Salary.

```
SELECT employee_id, last_name, salary,
       ROUND(salary * 1.155) AS "New Salary"
FROM employees;
```

3. Modify your query lab_03_02.sql to add a column that subtracts the old salary from the new salary. Label the column Increase.

```
SELECT employee_id, last_name, salary,
       ROUND(salary * 1.155) AS "New Salary",
       ROUND(salary * 0.155) AS "Increase"
FROM employees;
```

4. Write a query that displays the last name (with the first letter uppercase and all other letters lowercase) and the length of the last name for all employees whose name starts with the letters J, A, or M. Give each column an appropriate label. Sort the results by the employees' last names.

```
SELECT INITCAP(last_name) AS "Employee Name",
       LENGTH(last_name) AS "Name Length"
FROM employees
WHERE last_name LIKE 'J%' OR last_name LIKE 'A%' OR last_name LIKE 'M%'
ORDER BY last_name;
```

5. Rewrite the query so that the user is prompted to enter a letter that starts the last name. For example, if the user enters H when prompted for a letter, then the output should show all employees whose last name starts with the letter H.

```
User_input = 'H'
```

```
SELECT last_name FROM employees
WHERE last_name LIKE User_input || '%';
```

6. The HR department wants to find the length of employment for each employee. For each employee, display the last name and calculate the number of months between today and the date on which the employee was hired. Label the column MONTHS_WORKED. Order your results by the number of months employed. Round the number of months up to the closest whole number.

```
SELECT last_name,
       CEIL(MONTHS_BETWEEN(CURRENT_DATE, hire_date)) AS "MONTHS_WORKED"
FROM employees
ORDER BY "MONTHS_WORKED";
```

Note: Your results will differ.

7. Create a report that produces the following for each employee:
<employee last name> earns <salary> monthly but wants <3 times salary>. Label the column Dream Salaries.

```
SELECT last_name,
       salary AS "Current Salary",
       salary * 3 AS "Dream Salary"
FROM employees;
```

8. Create a query to display the last name and salary for all employees. Format the salary to be 15 characters long, left-padded with the \$ symbol. Label the column SALARY.

```
SELECT last_name,
       LPAD('$' || TO_CHAR(salary, '99999.99'), 15, ' ') AS "SALARY"
FROM employees;
```

9. Display each employee's last name, hire date, and salary review date, which is the first Monday after six months of service. Label the column REVIEW. Format the dates to appear in the format similar to "Monday, the Thirty-First of July, 2000."

```
SELECT last_name,
       hire_date,
       NEXT_DAY(ADD_MONTHS(hire_date, 6), 'MONDAY') AS "REVIEW"
FROM employees;
```

10. Display the last name, hire date, and day of the week on which the employee started. Label the column DAY. Order the results by the day of the week, starting with Monday.

```
SELECT last_name,  
       hire_date,  
       TO_CHAR(hire_date, 'DAY') AS "DAY"  
FROM employees  
ORDER BY TO_CHAR(hire_date, 'D');
```

Evaluation Procedure	Marks awarded
Query(5)	
Execution (5)	
Viva(5)	
Total (15)	
Faculty Signature	

EXERCISE-7

Displaying data from multiple tables

Objective

After the completion of this exercise, the students will be able to do the following:

- Write SELECT statements to access data from more than one table using equality and nonequality joins
- View data that generally does not meet a join condition by using outer joins
- Join a table to itself by using a self join

Sometimes you need to use data from more than one table.

Cartesian Products

- A Cartesian product is formed when:
 - A join condition is omitted
 - A join condition is invalid
 - All rows in the first table are joined to all rows in the second table
- To avoid a Cartesian product, always include a valid join condition in a WHERE clause.

A Cartesian product tends to generate a large number of rows, and the result is rarely useful. You should always include a valid join condition in a WHERE clause, unless you have a specific need to combine all rows from all tables.

Cartesian products are useful for some tests when you need to generate a large number of rows to simulate a reasonable amount of data.

Example:

To displays employee last name and department name from the EMPLOYEES and DEPARTMENTS tables.

```
SELECT last_name, department_name dept_name
FROM employees, departments;
```

Types of Joins

- Equijoin
- Non-equijoin
- Outer join
- Self join
- Cross joins
- Natural joins
- Using clause
- Full or two sided outer joins
- Arbitrary join conditions for outer joins

Joining Tables Using Oracle Syntax

```
SELECT table1.column, table2.column
FROM table1, table2
WHERE table1.column1 = table2.column2;
```

Write the join condition in the WHERE clause.

- Prefix the column name with the table name when the same column name appears in more than one table.

Guidelines

- When writing a SELECT statement that joins tables, precede the column name with the table name for clarity and to enhance database access.
- If the same column name appears in more than one table, the column name must be prefixed with the table name.
- To join n tables together, you need a minimum of n-1 join conditions. For example, to join four tables, a minimum of three joins is required. This rule may not apply if your table has a concatenated primary key, in which case more than one column is required to uniquely identify each row

What is an Equijoin?

To determine an employee's department name, you compare the value in the DEPARTMENT_ID column in the EMPLOYEES table with the DEPARTMENT_ID values in the DEPARTMENTS table.

The relationship between the EMPLOYEES and DEPARTMENTS tables is an equijoin—that is, values

in the DEPARTMENT_ID column on both tables must be equal. Frequently, this type of join involves

primary and foreign key complements.

Note: Equijoins are also called simple joins or inner joins

```
SELECT employees.employee_id, employees.last_name, employees.department_id,
       departments.department_id, departments.location_id
```

```
FROM   employees, departments
```

```
WHERE  employees.department_id = departments.department_id;
```

Additional Search Conditions

Using the AND Operator

Example:

To display employee Matos's department number and department name, you need an additional condition in the WHERE clause.

```
SELECT last_name, employees.department_id,
       department_name
```

```
FROM   employees, departments
```

```
WHERE  employees.department_id = departments.department_id AND   last_name = 'Matos';
```

Qualifying Ambiguous

Column Names

- Use table prefixes to qualify column names that are in multiple tables.
- Improve performance by using table prefixes.
- Distinguish columns that have identical names but reside in different tables by using column aliases.

Using Table Aliases

- Simplify queries by using table aliases.
- Improve performance by using table prefixes

Example:

```
SELECT e.employee_id, e.last_name, e.department_id,
       d.department_id, d.location_id
```

```
FROM   employees e , departments d
```

```
WHERE  e.department_id = d.department_id;
```

Joining More than Two Tables

To join n tables together, you need a minimum of n-1 join conditions. For example, to join three tables, a minimum of two joins is required.

Example:

To display the last name, the department name, and the city for each employee, you have to join the EMPLOYEES, DEPARTMENTS, and LOCATIONS tables.

```
SELECT e.last_name, d.department_name, l.city
FROM employees e, departments d, locations l
WHERE e.department_id = d.department_id
AND d.location_id = l.location_id;
```

Non-Equi Joins

A non-equijoin is a join condition containing something other than an equality operator. The relationship between the EMPLOYEES table and the JOB_GRADES table has an example of a non-equijoin. A relationship between the two tables is that the SALARY column in the EMPLOYEES table must be between the values in the LOWEST_SALARY and HIGHEST_SALARY columns of the JOB_GRADES table. The relationship is obtained using an operator other than equals (=).

Example:

```
SELECT e.last_name, e.salary, j.grade_level
FROM employees e, job_grades j
WHERE e.salary
BETWEEN j.lowest_sal AND j.highest_sal;
```

Outer Joins

Syntax

- You use an outer join to also see rows that do not meet the join condition.
- The Outer join operator is the plus sign (+).

```
SELECT table1.column, table2.column
FROM table1, table2
WHERE table1.column(+) = table2.column;
SELECT table1.column, table2.column
FROM table1, table2
WHERE table1.column = table2.column(+);
```

The missing rows can be returned if an outer join operator is used in the join condition. The operator is a plus sign enclosed in parentheses (+), and it is placed on the “side” of the join that is deficient in information. This operator has the effect of creating one or more null rows, to which one or more rows from the nondeficient table can be joined.

Example:

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e, departments d
WHERE e.department_id(+) = d.department_id ;
```

Outer Join Restrictions

- The outer join operator can appear on only one side of the expression—the side that has information missing. It returns those rows from one table that have no direct match in the other table.
- A condition involving an outer join cannot use the IN operator or be linked to another condition by the OR operator

Self Join

Sometimes you need to join a table to itself.

Example:

To find the name of each employee's manager, you need to join the EMPLOYEES table to itself, or perform a self join.

```
SELECT worker.last_name || ' works for '
|| manager.last_name
FROM employees worker, employees manager
WHERE worker.manager_id = manager.employee_id ;
```

Use a join to query data from more than one table.

```
SELECT table1.column, table2.column
FROM table1
[CROSS JOIN table2] |
[NATURAL JOIN table2] |
[JOIN table2 USING (column_name)] |
[JOIN table2
ON(table1.column_name = table2.column_name)] |
[LEFT|RIGHT|FULL OUTER JOIN table2
ON (table1.column_name = table2.column_name)];
```

In the syntax:

table1.column Denotes the table and column from which data is retrieved

CROSS JOIN Returns a Cartesian product from the two tables

NATURAL JOIN Joins two tables based on the same column name

JOIN table USING column_name Performs an equijoin based on the column name

JOIN table ON table1.column_name Performs an equijoin based on the condition in the ON clause
= table2.column_name

LEFT/RIGHT/FULL OUTER

Creating Cross Joins

- The CROSS JOIN clause produces the crossproduct of two tables.
- This is the same as a Cartesian product between the two tables.

Example:

```
SELECT last_name, department_name
FROM employees
CROSS JOIN departments ;
SELECT last_name, department_name
FROM employees, departments;
```

Creating Natural Joins

- The NATURAL JOIN clause is based on all columns in the two tables that have the same name.
- It selects rows from the two tables that have equal values in all matched columns.
- If the columns having the same names have different data types, an error is returned.

Example:

```
SELECT department_id, department_name,
location_id, city
FROM departments
NATURAL JOIN locations ;
```

LOCATIONS table is joined to the DEPARTMENT table by the LOCATION_ID column, which is the only column of the same name in both tables. If other common columns were present, the join would have used them all.

Example:

```
SELECT department_id, department_name,
location_id, city
FROM departments
NATURAL JOIN locations
WHERE department_id IN (20, 50);
```

Creating Joins with the USING Clause

- If several columns have the same names but the data types do not match, the NATURAL JOIN clause can be modified with the USING clause to specify the columns that should be used for an equijoin.
- Use the USING clause to match only one column when more than one column matches.
- Do not use a table name or alias in the referenced columns.
- The NATURAL JOIN and USING clauses are mutually exclusive.

Example:

```
SELECT l.city, d.department_name
FROM locations l JOIN departments d USING (location_id)
WHERE location_id = 1400;
EXAMPLE:
```

```
SELECT e.employee_id, e.last_name, d.location_id
FROM employees e JOIN departments d
USING (department_id) ;
```

Creating Joins with the ON Clause

- The join condition for the natural join is basically an equijoin of all columns with the same name.
- To specify arbitrary conditions or specify columns to join, the ON clause is used.
- The join condition is separated from other search conditions.
- The ON clause makes code easy to understand.

Example:

```
SELECT e.employee_id, e.last_name, e.department_id,
d.department_id, d.location_id
FROM employees e JOIN departments d
ON (e.department_id = d.department_id);
EXAMPLE:
```

```
SELECT e.last_name emp, m.last_name mgr
```



```
FROM employees e JOIN employees m
ON (e.manager_id = m.employee_id);
INNER Versus OUTER Joins
```

- A join between two tables that returns the results of the inner join as well as unmatched rows left (or right) tables is a left (or right) outer join.
- A join between two tables that returns the results of an inner join as well as the results of a left and right join is a full outer join.

LEFT OUTER JOIN

Example:

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e
LEFT OUTER JOIN departments d
ON (e.department_id = d.department_id) ;
```

Example of LEFT OUTER JOIN

This query retrieves all rows in the EMPLOYEES table, which is the left table even if there is no match in the DEPARTMENTS table.

This query was completed in earlier releases as follows:

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e, departments d
WHERE d.department_id (+) = e.department_id;
```

RIGHT OUTER JOIN

Example:

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e
RIGHT OUTER JOIN departments d
ON (e.department_id = d.department_id) ;
```

This query retrieves all rows in the DEPARTMENTS table, which is the right table even if there is no match in the EMPLOYEES table.

This query was completed in earlier releases as follows:

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e, departments d
WHERE d.department_id = e.department_id (+);
```

FULL OUTER JOIN

Example:

```
SELECT e.last_name, e.department_id, d.department_name
FROM employees e
FULL OUTER JOIN departments d
ON (e.department_id = d.department_id);
```

This query retrieves all rows in the EMPLOYEES table, even if there is no match in the DEPARTMENTS table. It also retrieves all rows in the DEPARTMENTS table, even if there is no match in the EMPLOYEES table.

Find the Solution for the following:

1. Write a query to display the last name, department number, and department name for all employees.

```
SELECT e.last_name, d.department_id, d.department_name
FROM employees e
JOIN departments d ON e.department_id = d.department_id;
```

2. Create a unique listing of all jobs that are in department 80. Include the location of the department in the output.

```
SELECT DISTINCT j.job_title, l.city, l.state_province
FROM jobs j
JOIN employees e ON j.job_id = e.job_id
JOIN departments d ON e.department_id = d.department_id
JOIN locations l ON d.location_id = l.location_id
WHERE d.department_id = 80;
```

3. Write a query to display the employee last name, department name, location ID, and city of all employees who earn a commission

```
SELECT e.last_name, d.department_name, l.location_id, l.city
FROM employees e
JOIN departments d ON e.department_id = d.department_id
JOIN locations l ON d.location_id = l.location_id
WHERE e.commission_pct IS NOT NULL;
```

2. Display the employee last name and department name for all employees who have an a(lowercase) in their last names. P

```
SELECT e.last_name, d.department_name
FROM employees e
```

```
JOIN departments d ON e.department_id = d.department_id  
WHERE e.last_name LIKE '%a%';
```

5. Write a query to display the last name, job, department number, and department name for all employees who work in Toronto.

```
SELECT e.last_name, j.job_title, e.department_id, d.department_name  
FROM employees e  
JOIN jobs j ON e.job_id = j.job_id  
JOIN departments d ON e.department_id = d.department_id  
JOIN locations l ON d.location_id = l.location_id  
WHERE l.city = 'Toronto';
```

6. Display the employee last name and employee number along with their manager's last name and manager number. Label the columns Employee, Emp#, Manager, and Mgr#, Respectively

```
SELECT e.last_name AS "Employee", e.employee_id AS "Emp#",  
       m.last_name AS "Manager", m.employee_id AS "Mgr#"  
FROM employees e  
LEFT JOIN employees m ON e.manager_id = m.employee_id;
```

7. Modify lab4_6.sql to display all employees including King, who has no manager. Order the results by the employee number.

```
SELECT e.last_name, e.employee_id, e.manager_id  
FROM employees e  
UNION  
SELECT 'King', 100, NULL  
FROM dual  
ORDER BY e.employee_id;
```

8. Create a query that displays employee last names, department numbers, and all the employees who work in the same department as a given employee. Give each column an appropriate label

```
SELECT e1.last_name, e1.department_id, e2.last_name AS "Colleague"  
FROM employees e1  
JOIN employees e2 ON e1.department_id = e2.department_id  
WHERE e1.employee_id = 101;
```

9. Show the structure of the JOB_GRADES table. Create a query that displays the name, job, department name, salary, and grade for all employees

```
CREATE TABLE job_grades (  
    grade CHAR(1),  
    lowest_sal NUMBER(8, 2) NOT NULL,  
    highest_sal NUMBER(8, 2) NOT NULL  
);
```

10. Create a query to display the name and hire date of any employee hired after employee Davies.

```
SELECT e.last_name, e.hire_date  
FROM employees e  
WHERE e.hire_date > (SELECT hire_date FROM employees WHERE last_name = 'Davies');
```

11. Display the names and hire dates for all employees who were hired before their managers, along with their manager's names and hire dates. Label the columns Employee, Emp Hired, Manager, and Mgr Hired, respectively.

```
SELECT e.last_name AS "Employee", e.hire_date AS "Emp Hired",  
m.last_name AS "Manager", m.hire_date AS "Mgr Hired"  
FROM employees e  
JOIN employees m ON e.manager_id = m.employee_id  
WHERE e.hire_date < m.hire_date;
```

Evaluation Procedure	Marks awarded
Query(5)	
Execution (5)	
Viva(5)	
Total (15)	
Faculty Signature	

EXERCISE-8

Aggregating Data Using Group Functions

Objectives

After the completion of this exercise, the students be will be able to do the following:

- Identify the available group functions
- Describe the use of group functions
- Group data by using the GROUP BY clause
- Include or exclude grouped rows by using the HAVING clause

What Are Group Functions?

Group functions operate on sets of rows to give one result per group

Types of Group Functions

- AVG
- COUNT
- MAX
- MIN
- STDDEV
- SUM
- VARIANCE

Each of the functions accepts an argument. The following table identifies the options that you can use in the syntax:

Function	Description
AVG ([DISTINCT <u>ALL</u>] <i>n</i>)	Average value of <i>n</i> , ignoring null values
COUNT ({ * [DISTINCT <u>ALL</u>] <i>expr</i> })	Number of rows, where <i>expr</i> evaluates to something other than null (count all selected rows using *, including duplicates and rows with nulls)
MAX ([DISTINCT <u>ALL</u>] <i>expr</i>)	Maximum value of <i>expr</i> , ignoring null values
MIN ([DISTINCT <u>ALL</u>] <i>expr</i>)	Minimum value of <i>expr</i> , ignoring null values
STDDEV ([DISTINCT <u>ALL</u>] <i>x</i>)	Standard deviation of <i>n</i> , ignoring null values
SUM ([DISTINCT <u>ALL</u>] <i>n</i>)	Sum values of <i>n</i> , ignoring null values
VARIANCE ([DISTINCT <u>ALL</u>] <i>x</i>)	Variance of <i>n</i> , ignoring null values

Group Functions: Syntax

```
SELECT [column,] group_function(column), ...  
FROM table  
[WHERE condition]  
[GROUP BY column]  
[ORDER BY column];
```

Guidelines for Using Group Functions

- DISTINCT makes the function consider only nonduplicate values; ALL makes it consider every value, including duplicates. The default is ALL and therefore does not need to be specified.

- The data types for the functions with an *expr* argument may be CHAR, VARCHAR2,

NUMBER, or DATE.

- All group functions ignore null values.

Using the AVG and SUM Functions

You can use AVG and SUM for numeric data.

```
SELECT AVG(salary), MAX(salary),  
MIN(salary), SUM(salary)  
FROM employees  
WHERE job_id LIKE '%REP%';
```

Using the MIN and MAX Functions

You can use MIN and MAX for numeric, character, and date data types.

```
SELECT MIN(hire_date), MAX(hire_date)  
FROM employees;
```

You can use the MAX and MIN functions for numeric, character, and date data types.
example displays the most junior and most senior employees.

The following example displays the employee last name that is first and the employee last name that is last in an alphabetized list of all employees:

```
SELECT MIN(last_name), MAX(last_name)  
FROM employees;
```

Note: The AVG, SUM, VARIANCE, and STDDEV functions can be used only with numeric data types. MAX and MIN cannot be used with LOB or LONG data types.

Using the COUNT Function

COUNT(*) returns the number of rows in a table:

```
SELECT COUNT(*)  
FROM employees  
WHERE department_id = 50;  
COUNT(expr) returns the number of rows with nonnull  
values for the expr:  
SELECT COUNT(commission_pct)  
FROM employees  
WHERE department_id = 80;
```

Using the DISTINCT Keyword

- COUNT(DISTINCT *expr*) returns the number of distinct non-null values of the *expr*.

- To display the number of distinct department values in the EMPLOYEES table:

```
SELECT COUNT(DISTINCT department_id) FROM employees;
```

Use the DISTINCT keyword to suppress the counting of any duplicate values in a column.

Group Functions and Null Values

Group functions ignore null values in the column:

```
SELECT AVG(commission_pct)
FROM employees;
```

The NVL function forces group functions to include null values:

```
SELECT AVG(NVL(commission_pct, 0))
FROM employees;
```

Creating Groups of Data

To divide the table of information into smaller groups. This can be done by using the GROUP BY clause.

GROUP BY Clause Syntax

```
SELECT column, group_function(column)
FROM table
[WHERE condition]
[GROUP BY group_by_expression]
[ORDER BY column];
```

In the syntax:

group_by_expression specifies columns whose values determine the basis for grouping rows

Guidelines

- If you include a group function in a SELECT clause, you cannot select individual results as well, *unless* the individual column appears in the GROUP BY clause. You receive an error message if you fail to include the column list in the GROUP BY clause.
- Using a WHERE clause, you can exclude rows before dividing them into groups.
- You must include the *columns* in the GROUP BY clause.
- You cannot use a column alias in the GROUP BY clause.

Using the GROUP BY Clause

All columns in the SELECT list that are not in group functions must be in the GROUP BY clause.

```
SELECT department_id, AVG(salary)
FROM employees
GROUP BY department_id ;
```

The GROUP BY column does not have to be in the SELECT list.

```
SELECT AVG(salary) FROM employees GROUP BY department_id ;
```

You can use the group function in the ORDER BY clause:

```
SELECT department_id, AVG(salary) FROM employees GROUP BY department_id ORDER BY  
AVG(salary);
```

Grouping by More Than One Column

```
SELECT department_id dept_id, job_id, SUM(salary) FROM employees  
GROUP BY department_id, job_id ;
```

Illegal Queries Using Group Functions

Any column or expression in the SELECT list that is not an aggregate function must be in the GROUP

BY clause:

```
SELECT department_id, COUNT(last_name) FROM employees;
```

You can correct the error by adding the GROUP BY clause:

```
SELECT department_id, count(last_name) FROM employees GROUP BY department_id;
```

You cannot use the WHERE clause to restrict groups.

- You use the HAVING clause to restrict groups.
- You cannot use group functions in the WHERE clause.

```
SELECT department_id, AVG(salary) FROM employees WHERE AVG(salary) > 8000  
GROUP BY department_id;
```

You can correct the error in the example by using the HAVING clause to restrict groups:

```
SELECT department_id, AVG(salary) FROM employees  
HAVING AVG(salary) > 8000 GROUP BY department_id;
```

Restricting Group Results

With the HAVING Clause .When you use the HAVING clause, the Oracle server restricts groups as follows:

1. Rows are grouped.
2. The group function is applied.
3. Groups matching the HAVING clause are displayed.

Using the HAVING Clause

```
SELECT department_id, MAX(salary) FROM employees  
GROUP BY department_id HAVING MAX(salary)>10000 ;
```

The following example displays the department numbers and average salaries for those departments with a maximum salary that is greater than \$10,000:

```
SELECT department_id, AVG(salary) FROM employees GROUP BY department_id  
HAVING max(salary)>10000;
```


Example displays the job ID and total monthly salary for each job that has a total payroll exceeding \$13,000. The example excludes sales representatives and sorts the list by the total monthly salary.

```
SELECT job_id, SUM(salary) PAYROLL FROM employees WHERE job_id NOT LIKE  
'%REP%'  
GROUP BY job_id HAVING SUM(salary) > 13000 ORDER BY SUM(salary);
```

Nesting Group Functions

Display the maximum average salary:

Group functions can be nested to a depth of two. The slide example displays the maximum average salary.

```
SELECT MAX(AVG(salary)) FROM employees GROUP BY department_id;
```

Summary

In this exercise, students should have learned how to:

- Use the group functions COUNT, MAX, MIN, and AVG
- Write queries that use the GROUP BY clause
- Write queries that use the HAVING clause

```
SELECT column, group_function  
FROM table  
[WHERE condition]  
[GROUP BY group_by_expression]  
[HAVING group_condition]  
[ORDER BY column];
```

Find the Solution for the following:

Determine the validity of the following three statements. Circle either True or False.

1. Group functions work across many rows to produce one result per group.
True/False -- True
2. Group functions include nulls in calculations.
True/False -- False
3. The WHERE clause restricts rows prior to inclusion in a group calculation.
True/False --False

The HR department needs the following reports:

4. Find the highest, lowest, sum, and average salary of all employees. Label the columns Maximum, Minimum, Sum, and Average, respectively. Round your results to the nearest whole number

```
SELECT  
    ROUND(MAX(salary)) AS Maximum,  
    ROUND(MIN(salary)) AS Minimum,  
    ROUND(SUM(salary)) AS Sum,  
    ROUND(AVG(salary)) AS Average  
FROM employees;
```

5. Modify the above query to display the minimum, maximum, sum, and average salary for each job type.

```
SELECT job_id,  
       ROUND(MIN(salary)) AS Minimum,  
       ROUND(MAX(salary)) AS Maximum,  
       ROUND(SUM(salary)) AS Sum,  
       ROUND(AVG(salary)) AS Average  
FROM employees  
GROUP BY job_id;
```

6. Write a query to display the number of people with the same job. Generalize the query so that the user in the HR department is prompted for a job title.

```
SELECT job_id, COUNT(*) AS "Number of People"  
FROM employees  
WHERE job_id = :user_input_job_id  
GROUP BY job_id;
```

7. Determine the number of managers without listing them. Label the column Number of Managers. *Hint: Use the MANAGER_ID column to determine the number of managers.*

```
SELECT COUNT(DISTINCT manager_id) AS "Number of Managers"  
FROM employees;
```

8. Find the difference between the highest and lowest salaries. Label the column DIFFERENCE.

```
SELECT ROUND(MAX(salary) - MIN(salary)) AS Difference  
FROM employees;
```

9. Create a report to display the manager number and the salary of the lowest-paid employee for that manager. Exclude anyone whose manager is not known. Exclude any groups where the minimum salary is \$6,000 or less. Sort the output in descending order of salary.

```
SELECT manager_id, MIN(salary) AS "Lowest Salary"
FROM employees
WHERE manager_id IS NOT NULL
GROUP BY manager_id
HAVING MIN(salary) > 6000
ORDER BY "Lowest Salary" DESC;
```

10. Create a query to display the total number of employees and, of that total, the number of employees hired in 1995, 1996, 1997, and 1998. Create appropriate column headings.

```
SELECT
    COUNT(*) AS "Total Employees",
    SUM(CASE WHEN EXTRACT(YEAR FROM hire_date) = 1995 THEN 1 ELSE 0 END) AS
"Hired in 1995",
    SUM(CASE WHEN EXTRACT(YEAR FROM hire_date) = 1996 THEN 1 ELSE 0 END) AS
"Hired in 1996",
    SUM(CASE WHEN EXTRACT(YEAR FROM hire_date) = 1997 THEN 1 ELSE 0 END) AS
"Hired in 1997",
    SUM(CASE WHEN EXTRACT(YEAR FROM hire_date) = 1998 THEN 1 ELSE 0 END) AS
"Hired in 1998"
FROM employees;
```

11. Create a matrix query to display the job, the salary for that job based on department number, and the total salary for that job, for departments 20, 50, 80, and 90, giving each column an appropriate heading.

```
SELECT
    j.job_title AS "Job",
    ROUND(AVG(CASE WHEN d.department_id = 20 THEN e.salary ELSE NULL END), 2) AS
"Salary in Dept 20",
    ROUND(AVG(CASE WHEN d.department_id = 50 THEN e.salary ELSE NULL END), 2) AS
"Salary in Dept 50",
    ROUND(AVG(CASE WHEN d.department_id = 80 THEN e.salary ELSE NULL END), 2) AS
"Salary in Dept 80",
    ROUND(AVG(CASE WHEN d.department_id = 90 THEN e.salary ELSE NULL END), 2) AS
"Salary in Dept 90",
    ROUND(SUM(e.salary), 2) AS "Total Salary"
FROM employees e
JOIN jobs j ON e.job_id = j.job_id
JOIN departments d ON e.department_id = d.department_id
WHERE d.department_id IN (20, 50, 80, 90)
GROUP BY j.job_title;
```

12. Write a query to display each department's name, location, number of employees, and the average salary for all the employees in that department. Label the column name-Location, Number of people, and salary respectively. Round the average salary to two decimal places.

```
SELECT
    d.department_name AS "Department Name",
    l.city || ', ' || l.state_province AS "Location",
    COUNT(*) AS "Number of People",
    ROUND(AVG(e.salary), 2) AS "Salary"
FROM employees e
JOIN departments d ON e.department_id = d.department_id
JOIN locations l ON d.location_id = l.location_id
GROUP BY d.department_name, l.city, l.state_province;
```

Evaluation Procedure	Marks awarded
Query(5)	
Execution (5)	
Viva(5)	
Total (15)	
Faculty Signature	

EXERCISE-9

Sub queries

Objectives

After completing this lesson, you should be able to do the following:

- Define subqueries
- Describe the types of problems that subqueries can solve
- List the types of subqueries
- Write single-row and multiple-row subqueries

Using a Subquery to Solve a Problem

Who has a salary greater than Abel's?

Main query:

Which employees have salaries greater than Abel's salary?

Subquery:

What is Abel's salary?

Subquery Syntax

`SELECT select_list FROM table WHERE expr operator (SELECT select_list FROM table);`

- The subquery (inner query) executes once before the main query (outer query).
- The result of the subquery is used by the main query.

A subquery is a SELECT statement that is embedded in a clause of another SELECT statement. You can build powerful statements out of simple ones by using subqueries. They can be very useful when you need to select rows from a table with a condition that depends on the data in the table itself.

You can place the subquery in a number of SQL clauses, including the following:

- WHERE clause
- HAVING clause
- FROM clause

In the syntax:

operator includes a comparison condition such as >, =, or IN

Note: Comparison conditions fall into two classes: single-row operators

(>, =, >=, <, <=, <>) and multiple-row operators (IN, ANY, ALL). statement. The subquery generally executes first, and its output is used to complete the query condition for the main (or outer) query

Using a Subquery

`SELECT last_name FROM employees WHERE salary > (SELECT salary FROM employees WHERE last_name = 'Abel');`

The inner query determines the salary of employee Abel. The outer query takes the result of the inner query and uses this result to display all the employees who earn more than this amount.

Guidelines for Using Subqueries

- Enclose subqueries in parentheses.
- Place subqueries on the right side of the comparison condition.
- The ORDER BY clause in the subquery is not needed unless you are performing Top-N analysis.
- Use single-row operators with single-row

subqueries, and use multiple-row operators with multiple-row subqueries.

Types of Subqueries

- Single-row subqueries: Queries that return only one row from the inner SELECT statement.
- Multiple-row subqueries: Queries that return more than one row from the inner SELECT statement.

Single-Row Subqueries

- Return only one row
- Use single-row comparison operators

Example

Display the employees whose job ID is the same as that of employee 141:

```
SELECT last_name, job_id FROM employees WHERE job_id = (SELECT job_id FROM
employees
WHERE employee_id = 141);
```

Displays employees whose job ID is the same as that of employee 141 and whose salary is greater than that of employee 143.

```
SELECT last_name, job_id, salary FROM employees WHERE job_id =(SELECT job_id FROM
employees WHERE employee_id = 141) AND salary > (SELECT salary FROM employees
WHERE employee_id = 143);
```

Using Group Functions in a Subquery

Displays the employee last name, job ID, and salary of all employees whose salary is equal to the minimum salary. The MIN group function returns a single value (2500) to the outer query.

```
SELECT last_name, job_id, salary FROM employees WHERE salary = (SELECT MIN(salary)
FROM employees);
```

The HAVING Clause with Subqueries

- The Oracle server executes subqueries first.
- The Oracle server returns results into the HAVING clause of the main query.

Displays all the departments that have a minimum salary greater than that of department 50.

```
SELECT department_id, MIN(salary)
FROM employees
GROUP BY department_id
HAVING MIN(salary) >
(SELECT MIN(salary)
FROM employees
WHERE department_id = 50);
```

Example

Find the job with the lowest average salary.

```
SELECT job_id, AVG(salary)
FROM employees
GROUP BY job_id
HAVING AVG(salary) = (SELECT MIN(AVG(salary))
FROM employees
GROUP BY job_id);
```

What Is Wrong in this Statements?

```
SELECT employee_id, last_name
FROM employees
WHERE salary =(SELECT MIN(salary) FROM employees GROUP BY department_id);
Will This Statement Return Rows?
SELECT last_name, job_id
FROM employees
WHERE job_id =(SELECT job_id FROM employees WHERE last_name = 'Haas');
```

Multiple-Row Subqueries

- Return more than one row
- Use multiple-row comparison operators

Example

Find the employees who earn the same salary as the minimum salary for each department.

```
SELECT last_name, salary, department_id FROM employees WHERE salary IN (SELECT
MIN(salary)
FROM employees GROUP BY department_id);
```

Using the ANY Operator in Multiple-Row Subqueries

```
SELECT employee_id, last_name, job_id, salary FROM employees WHERE salary < ANY
(SELECT salary FROM employees WHERE job_id = 'IT_PROG') AND job_id <> 'IT_PROG';
```

Displays employees who are not IT programmers and whose salary is less than that of any IT programmer. The maximum salary that a programmer earns is \$9,000.

< ANY means less than the maximum. >ANY means more than the minimum. =ANY is equivalent to IN.

Using the ALL Operator in Multiple-Row Subqueries

```
SELECT employee_id, last_name, job_id, salary
FROM employees
WHERE salary < ALL (SELECT salary FROM employees WHERE job_id = 'IT_PROG')
AND job_id <> 'IT_PROG';
```

Displays employees whose salary is less than the salary of all employees with a job ID of IT_PROG and whose job is not IT_PROG.

➤ ALL means more than the maximum, and <ALL means less than the minimum.

The NOT operator can be used with IN, ANY, and ALL operators.

Null Values in a Subquery

```
SELECT emp.last_name FROM employees emp
WHERE emp.employee_id NOT IN (SELECT mgr.manager_id FROM employees mgr);
```

Notice that the null value as part of the results set of a subquery is not a problem if you use the IN operator. The IN operator is equivalent to =ANY. For example, to display the employees who have subordinates, use the following SQL statement:

```
SELECT emp.last_name
FROM employees emp
WHERE emp.employee_id IN (SELECT mgr.manager_id FROM employees mgr);
```

Display all employees who do not have any subordinates:

```
SELECT last_name FROM employees
WHERE employee_id NOT IN (SELECT manager_id FROM employees WHERE manager_id IS
NOT NULL);
```

Find the Solution for the following:

1. The HR department needs a query that prompts the user for an employee last name. The query then displays the last name and hire date of any employee in the same department as the employee whose name they supply (excluding that employee). For example, if the user enters Zlotkey, find all employees who work with Zlotkey (excluding Zlotkey).

```
SELECT last_name, hire_date
FROM employees
WHERE department_id = (SELECT department_id FROM employees WHERE last_name =
'&name')
AND last_name <> 'Zlotkey';
```

2. Create a report that displays the employee number, last name, and salary of all employees who earn more than the average salary. Sort the results in order of ascending salary.


```
SELECT employee_id, last_name, salary
FROM employees
WHERE salary > (SELECT AVG(salary) FROM employees);
```

3. Write a query that displays the employee number and last name of all employees who work in a department with any employee whose last name contains a *u*.

```
SELECT employee_id, last_name
FROM employees
WHERE department_id IN
  (SELECT department_id FROM employees WHERE last_name LIKE '%u%');
```

4. The HR department needs a report that displays the last name, department number, and job ID of all employees whose department location ID is 1700.

```
SELECT last_name, department_id, job_id
FROM employees
WHERE department_id IN
  (SELECT department_id FROM departments WHERE location_id = 1700);
```

5. Create a report for HR that displays the last name and salary of every employee who reports to King.

```
SELECT last_name, salary
FROM employees
WHERE manager_id = (SELECT employee_id FROM employees WHERE last_name = 'King');
```

6. Create a report for HR that displays the department number, last name, and job ID for every employee in the Executive department.

```
SELECT department_id, last_name, job_id
FROM employees
WHERE department_id = (SELECT department_id FROM departments WHERE
  department_name = 'Executive');
```

7. Modify the query 3 to display the employee number, last name, and salary of all employees who earn more than the average salary and who work in a department with any employee whose last name contains a *u*.

```
SELECT e.employee_id, e.last_name, e.salary
FROM employees e
WHERE e.department_id IN
      (SELECT d.department_id FROM departments d WHERE d.location_id = 1700)
AND e.salary > (SELECT AVG(salary) FROM employees);
```

Evaluation Procedure	Marks awarded
Query(5)	
Execution (5)	
Viva(5)	
Total (15)	
Faculty Signature	

EXERCISE-10

USING THE SET OPERATORS

Objectives

After the completion this exercise, the students should be able to do the following:

- Describe set operators
- Use a set operator to combine multiple queries into a single query
- Control the order of rows returned

The set operators combine the results of two or more component queries into one result.

Queries containing set operators are called *compound queries*.

Operator	Returns
UNION	All distinct rows selected by either query
UNION ALL	All rows selected by either query, including all duplicates
INTERSECT	All distinct rows selected by both queries
MINUS	All distinct rows that are selected by the first SELECT statement and not selected in the second SELECT statement

The tables used in this lesson are:

- EMPLOYEES: Provides details regarding all current employees
- JOB_HISTORY: Records the details of the start date and end date of the former job, and the job identification number and department when an employee switches jobs

UNION Operator

Guidelines

- The number of columns and the data types of the columns being selected must be identical in all the SELECT statements used in the query. The names of the columns need not be identical.
- UNION operates over all of the columns being selected.
- NULL values are not ignored during duplicate checking.
- The IN operator has a higher precedence than the UNION operator.
- By default, the output is sorted in ascending order of the first column of the SELECT clause.

Example:

Display the current and previous job details of all employees. Display each employee only once.

```
SELECT employee_id, job_id FROM employees UNION SELECT employee_id, job_id  
FROM job_history;
```

Example:

```
SELECT employee_id, job_id, department_id
FROM employees
UNION
SELECT employee_id, job_id, department_id
FROM job_history;
```

UNION ALL Operator

Guidelines

The guidelines for UNION and UNION ALL are the same, with the following two exceptions that pertain to UNION ALL:

- Unlike UNION, duplicate rows are not eliminated and the output is not sorted by default.
- The DISTINCT keyword cannot be used.

Example:

Display the current and previous departments of all employees.

```
SELECT employee_id, job_id, department_id
FROM employees
UNION ALL
SELECT employee_id, job_id, department_id
FROM job_history
ORDER BY employee_id;
```

INTERSECT Operator

Guidelines

- The number of columns and the data types of the columns being selected by the SELECT statements in the queries must be identical in all the SELECT statements used in the query. The names of the columns need not be identical.
- Reversing the order of the intersected tables does not alter the result.
- INTERSECT does not ignore NULL values.

Example:

Display the employee IDs and job IDs of those employees who currently have a job title that is the same as their job title when they were initially hired (that is, they changed jobs but have now gone back to doing their original job).

```
SELECT employee_id, job_id FROM employees
INTERSECT
SELECT employee_id, job_id
FROM job_history;
```

Example

```
SELECT employee_id, job_id, department_id
FROM employees
INTERSECT
SELECT employee_id, job_id, department_id
FROM job_history;
```

MINUS Operator

Guidelines

- The number of columns and the data types of the columns being selected by the SELECT statements in the queries must be identical in all the SELECT statements used in the query. The names of the columns need not be identical.
- All of the columns in the WHERE clause must be in the SELECT clause for the MINUS operator to work.

Example:

Display the employee IDs of those employees who have not changed their jobs even once.

```
SELECT employee_id, job_id
FROM employees
MINUS
SELECT employee_id, job_id
FROM job_history;
```

Find the Solution for the following:

1. The HR department needs a list of department IDs for departments that do not contain the job ID ST_CLERK. Use set operators to create this report.

```
SELECT DISTINCT department_id
FROM employees
WHERE department_id NOT IN (SELECT department_id FROM employees WHERE job_id =
'ST_CLERK');
```

2. The HR department needs a list of countries that have no departments located in them. Display the country ID and the name of the countries. Use set operators to create this report.

```
SELECT country_id, country_name
FROM countries
WHERE country_id NOT IN
(SELECT DISTINCT location_id FROM departments);
```

3. Produce a list of jobs for departments 10, 50, and 20, in that order. Display job ID and department ID using set operators.

```
SELECT job_id, department_id
FROM employees
WHERE department_id = 10
UNION ALL
SELECT job_id, department_id
FROM employees
WHERE department_id = 50
UNION ALL
SELECT job_id, department_id
FROM employees
WHERE department_id = 20;
```

4. Create a report that lists the employee IDs and job IDs of those employees who currently have a job title that is the same as their job title when they were initially hired by the company (that is, they changed jobs but have now gone back to doing their original job).

```
SELECT e.employee_id, e.job_id
FROM employees e
WHERE e.job_id IN
      (SELECT j.job_id FROM job_history j WHERE j.start_date = e.hire_date);
```

5. The HR department needs a report with the following specifications:

- Last name and department ID of all the employees from the EMPLOYEES table, regardless of whether or not they belong to a department.

- Department ID and department name of all the departments from the DEPARTMENTS table, regardless of whether or not they have employees working in them Write a compound query to accomplish this.

```
SELECT last_name, department_id FROM employees
UNION ALL
SELECT NULL, department_id, department_name FROM departments;
```

Evaluation Procedure	Marks awarded
Query(5)	
Execution (5)	
Viva(5)	
Total (15)	
Faculty Signature	

EXERCISE-11

CREATING VIEWS

After the completion of this exercise, students will be able to do the following:

- Describe a view
- Create, alter the definition of, and drop a view
- Retrieve data through a view
- Insert, update, and delete data through a view
- Create and use an inline view

View

A view is a logical table based on a table or another view. A view contains no data but is like a window through which data from tables can be viewed or changed. The tables on which a view is based are called base tables.

Advantages of Views

- To restrict data access
- To make complex queries easy
- To provide data independence
- To present different views of the same data

Classification of views

1. Simple view
2. Complex view

Feature	Simple	Complex
No. of tables	One	One or more
Contains functions	No	Yes
Contains groups of data	No	Yes
DML operations thr' view	Yes	Not always

Creating a view

Syntax

CREATE OR REPLACE FORCE/NOFORCE VIEW view_name AS Subquery WITH CHECK OPTION CONSTRAINT constraint WITH READ ONLY CONSTRAINT constraint;

FORCE - Creates the view regardless of whether or not the base tables exist.

NOFORCE - Creates the view only if the base table exist.

WITH CHECK OPTION CONSTRAINT-specifies that only rows accessible to the view can be inserted or updated.

WITH READ ONLY CONSTRAINT-ensures that no DML operations can be performed on the view.

Example: 1 (Without using Column aliases)

Create a view EMPVU80 that contains details of employees in department80.

Example 2:

```
CREATE VIEW empvu80 AS SELECT employee_id, last_name, salary FROM employees
WHERE department_id=80;
```

Example:1 (Using column aliases)

```
CREATE VIEW salvu50
AS SELECT employee_id,id_number, last_name NAME, salary *12 ANN_SALARY
FROM employees
WHERE department_id=50;
```

Retrieving data from a view**Example:**

```
SELECT * from salvu50;
```

Modifying a view

A view can be altered without dropping, re-creating.

Example: (Simple view)

Modify the EMPVU80 view by using CREATE OR REPLACE.

```
CREATE OR REPLACE VIEW empvu80 (id_number, name, sal, department_id)
AS SELECT employee_id,first_name, last_name, salary, department_id
FROM employees
WHERE department_id=80;
```

Example: (complex view)

```
CREATE VIEW dept_sum_vu (name, minsal, maxsal,avgsal)
AS SELECT d.department_name, MIN(e.salary), MAX(e.salary), AVG(e.salary)
FROM employees e, department d
WHERE e.deparment_id=d.deparment_id
GROUP BY d.department_name;
```

Rules for performing DML operations on view

- Can perform operations on simple views
- Cannot remove a row if the view contains the following:
 - Group functions
 - Group By clause
 - Distinct keyword
- Cannot modify data in a view if it contains
 - Group functions
 - Group By clause
 - Distinct keyword
 - Columns contain by expressions

-
- Cannot add data thr' a view if it contains
 - Group functions
 - Group By clause
 - Distinct keyword
 - Columns contain by expressions
 - NOT NULL columns in the base table that are not selected by the view

Example: (Using the WITH CHECK OPTION clause)

```
CREATE OR REPLACE VIEW empvu20
AS SELECT *
FROM employees
WHERE department_id=20
WITH CHECK OPTION CONSTRAINT empvu20_ck;
```

Note:Any attempt to change the department number for any row in the view fails because it violates the WITH CHECK OPTION constraint.

Example – (Execute this and note the error)

```
UPDATE empvu20 SET department_id=10 WHERE employee_id=201;
```

Denying DML operations

Use of WITH READ ONLY option.

Any attempt to perform a DML on any row in the view results in an oracle server error.

Try this code:

```
CREATE OR REPLACE VIEW empvu10(employee_number, employee_name,job_title)
AS SELECT employee_id, last_name, job_id
FROM employees
WHERE department_id=10
WITH READ ONLY;
```

Find the Solution for the following:

1. Create a view called EMPLOYEE_VU based on the employee numbers, employee names and department numbers from the EMPLOYEES table. Change the heading for the employee name to EMPLOYEE.

```
CREATE OR REPLACE VIEW EMPLOYEE_VU AS
SELECT employee_id, last_name AS employee, department_id
FROM employees;
```

2. Display the contents of the EMPLOYEES_VU view.

```
SELECT * FROM EMPLOYEE_VU;
```

3. Select the view name and text from the USER_VIEWS data dictionary views.

```
SELECT view_name, text  
FROM user_views;
```

4. Using your EMPLOYEES_VU view, enter a query to display all employees names and department.

```
SELECT employee, department_id  
FROM EMPLOYEE_VU;
```

5. Create a view named DEPT50 that contains the employee number, employee last names and department numbers for all employees in department 50. Label the view columns EMPNO, EMPLOYEE and DEPTNO. Do not allow an employee to be reassigned to another department through the view.

```
CREATE OR REPLACE VIEW DEPT50 AS  
SELECT employee_id AS EMPNO, last_name AS EMPLOYEE, department_id AS DEPTNO  
FROM employees  
WHERE department_id = 50;
```

6. Display the structure and contents of the DEPT50 view.

```
DESCRIBE DEPT50;  
SELECT * FROM DEPT50;
```

7. Attempt to reassign Matos to department 80.

```
UPDATE DEPT50 SET DEPTNO = 80 WHERE EMPLOYEE = 'Matos';
```

8. Create a view called SALARY_VU based on the employee last names, department names, salaries, and salary grades for all employees. Use the Employees, DEPARTMENTS and JOB_GRADE tables. Label the column Employee, Department, salary, and Grade respectively.

```
CREATE OR REPLACE VIEW SALARY_VU AS
SELECT e.last_name AS Employee, d.department_name AS Department,
       e.salary, j.grade
FROM employees e
JOIN departments d ON e.department_id = d.department_id
JOIN job_grade j ON e.salary BETWEEN j.lowest_sal AND j.highest_sal;
```

Evaluation Procedure	Marks awarded
Query(5)	
Execution (5)	
Viva(5)	
Total (15)	
Faculty Signature	

EXERCISE 12

Intro to Constraints; NOT NULL and UNIQUE Constraints

Global Fast Foods has been very successful this past year and has opened several new stores. They need to add a table to their database to store information about each of their store's locations. The owners want to make sure that all entries have an identification number, date opened, address, and city and that no other entry in the table can have the same email address. Based on this information, answer the following questions about the global_locations table. Use the table for your answers.

Global Fast Foods global_locations Table						
NAME	TYPE	LENGTH	PRECISION	SCALE	NULLABLE	DEFAULT
Id						
name						
date_opened						
address						
city						
zip/postal code						
phone						
email						
manager_id						
Emergency contact						

1. What is a “constraint” as it relates to data integrity?

A constraint in the context of data integrity refers to predefined rules or conditions that are enforced on a database to ensure the accuracy and consistency of data

2. What are the limitations of constraints that may be applied at the column level and at the table level?

- Constraints can be applied at both the **column level** and the **table level**. Here are the key differences:

-

Column-Level Constraints:

- Apply only to a specific column.
- Examples include **NOT NULL**, **UNIQUE**, and **CHECK** constraints.
- Cannot involve multiple columns.

-

Table-Level Constraints:

- Apply to the entire table.
- Can involve multiple columns.
- Examples include **PRIMARY KEY**, **FOREIGN KEY**, and **CHECK** constraints.

3. Why is it important to give meaningful names to constraints?

For clarity, maintenance, documentation and debugging.

4. Based on the information provided by the owners, choose a datatype for each column. Indicate the length, precision, and scale for each NUMBER datatype.

Id: NUMBER (10,0)
name: VARCHAR2 (100) (nullable)
date_opened: DATE (nullable)
address: VARCHAR2 (200) (nullable)
city: VARCHAR2 (100) (nullable)
zip/postal code: VARCHAR2 (20) (nullable)
phone: VARCHAR2 (20) (nullable)
email: VARCHAR2 (100)
manager_id: NUMBER (10,0) (nullable)
Emergency contact: VARCHAR2 (100) (nullable)

5. Use "(nullable)" to indicate those columns that can have null values.

name (nullable)
date_opened (nullable)
address (nullable)
city (nullable)
zip/postal code (nullable)
phone (nullable)
manager_id (nullable)
Emergency contact (nullable)

6. Write the CREATE TABLE statement for the Global Fast Foods locations table to define the constraints at the column level.

```
CREATE TABLE global_locations (  
  Id NUMBER(10,0) PRIMARY KEY,  
  name VARCHAR2(100),  
  date_opened DATE,  
  address VARCHAR2(200),  
  city VARCHAR2(100),  
  "zip/postal code" VARCHAR2(20),  
  phone VARCHAR2(20),  
  email VARCHAR2(100) UNIQUE,  
  manager_id NUMBER(10,0),  
  "Emergency contact" VARCHAR2(100)  
);
```

7. Execute the CREATE TABLE statement in Oracle Application Express.

```
CREATE TABLE global_locations (  
  Id NUMBER(10,0) PRIMARY KEY,  
  name VARCHAR2(100),  
  date_opened DATE,  
  address VARCHAR2(200),  
  city VARCHAR2(100),  
  "zip/postal code" VARCHAR2(20),  
  phone VARCHAR2(20),  
  email VARCHAR2(100) UNIQUE,  
  manager_id NUMBER(10,0),  
  "Emergency contact" VARCHAR2(100)  
);
```

8. Execute a DESCRIBE command to view the Table Summary information.

```
DESCRIBE global_locations;
```

9. Rewrite the CREATE TABLE statement for the Global Fast Foods locations table to define the UNIQUE constraints at the table level. Do not execute this statement.

NAME	TYPE	LENGTH	PRECISION	SCALE	NULLABLE	DEFAULT
id	number	4				
loc_name	varchar2	20			X	
	date					
address	varchar2	30				
city	varchar2	20				
zip_postal	varchar2	20			X	
phone	varchar2	15			X	
email	varchar2	80			X	
manager_id	number	4			X	
contact	varchar2	40			X	

```
CREATE TABLE global_locations (  
  Id NUMBER(10,0) PRIMARY KEY,  
  name VARCHAR2(100),  
  date_opened DATE,  
  address VARCHAR2(200),  
  city VARCHAR2(100),  
  "zip/postal code" VARCHAR2(20),  
  phone VARCHAR2(20),  
  email VARCHAR2(100),  
  manager_id NUMBER(10,0),  
  "Emergency contact" VARCHAR2(100),  
  CONSTRAINT unique_email UNIQUE (email)  
);
```

PRIMARY KEY, FOREIGN KEY, and CHECK Constraints

1. What is the purpose of a
 - PRIMARY KEY
 - FOREIGN KEY
 - CHECK CONSTRAINT

PRIMARY KEY: Ensures each record in a table is uniquely identifiable and serves as a unique identifier for each row. It also automatically creates a unique index on the primary key column(s).

FOREIGN KEY: Maintains referential integrity between two related tables. It establishes a link between a column in one table (child table) and a column in another table (parent table), ensuring that the values in the child table exist in the parent table.

CHECK CONSTRAINT: Validates the values entered into a column to ensure they meet a specific condition or set of conditions. It restricts the range of values that can be entered into a column.

2. Using the column information for the animals table below, name constraints where applicable at the table level, otherwise name them at the column level. Define the primary key (animal_id). The license_tag_number must be unique. The admit_date and vaccination_date columns cannot contain null values.

```
animal_id NUMBER(6)
name VARCHAR2(25)
license_tag_number NUMBER(10)
admit_date DATE
adoption_id NUMBER(5),
vaccination_date DATE
```

- **PRIMARY KEY:** animal_id (column-level)
- **UNIQUE CONSTRAINT:** license_tag_number (column-level)
- **NOT NULL CONSTRAINT:** admit_date, vaccination_date (column-level)

3. Create the animals table. Write the syntax you will use to create the table.

```
CREATE TABLE animals (
  animal_id NUMBER(6) PRIMARY KEY,
  name VARCHAR2(25),
  license_tag_number NUMBER(10) UNIQUE,
  admit_date DATE NOT NULL,
  adoption_id NUMBER(5),
  vaccination_date DATE NOT NULL
);
```

4. Enter one row into the table. Execute a SELECT * statement to verify your input. Refer to the graphic below for input.

ANIMAL_ID	NAME	LICENSE_TAG_NUMBER	ADMIT_DATE	ADOPTION_ID	VACCINATION_DATE
101	Spot	35540	10-Oct-2004	205	12-Oct-2004

```
INSERT INTO animals (animal_id, name, license_tag_number, admit_date, adoption_id,
vaccination_date)
VALUES (101, 'Spot', 35540, TO_DATE('10-Oct-2004', 'DD-Mon-YYYY'), 205, TO_DATE('12-Oct-2004',
'DD-Mon-YYYY'));
```

```
SELECT * FROM animals;
```

5. Write the syntax to create a foreign key (adoption_id) in the animals table that has a corresponding primary-key reference in the adoptions table. Show both the column-level and table-level syntax. Note that because you have not actually created an adoptions table, no adoption_id primary key exists, so the foreign key cannot be added to the animals table.

```
ALTER TABLE animals
ADD CONSTRAINT fk_adoption_id FOREIGN KEY (adoption_id) REFERENCES adoptions(adoption_id);
```

```
ALTER TABLE animals
ADD CONSTRAINT fk_adoption_id FOREIGN KEY (adoption_id)
REFERENCES adoptions(adoption_id);
```

6. What is the effect of setting the foreign key in the ANIMAL table as:

- a. ON DELETE CASCADE
- b. ON DELETE SET NULL

ON DELETE CASCADE: If a record in the parent table (adoptions) is deleted, all corresponding records in the child table (animals) will be automatically deleted.

ON DELETE SET NULL: If a record in the parent table (adoptions) is deleted, the corresponding foreign key column values in the child table (animals) will be set to NULL.

7. What are the restrictions on defining a CHECK constraint?

The condition in the CHECK constraint must be a Boolean expression evaluating to true or false.
The condition cannot reference other tables; it must apply only to the column(s) within the same table.
Certain complex expressions might not be allowed, depending on the database system.

Evaluation Procedure	Marks awarded
Query(5)	
Execution (5)	
Viva(5)	
Total (15)	
Faculty Signature	

EXERCISE 13

Creating Views

1. What are three uses for a view from a DBA's perspective?

Security: Views can restrict access to sensitive data by showing only certain columns or rows to users based on their permissions.

Simplification: Views can simplify complex queries by predefining joins, aggregations, or calculations, making it easier for users to query data without needing to understand the underlying database schema.

Performance: Views can improve query performance by storing the results of complex queries and allowing users to retrieve them quickly without re-executing the underlying query logic.

2. Create a simple view called view_d_songs that contains the ID, title and artist from the DJs on Demand table for each "New Age" type code. In the subquery, use the alias "Song Title" for the title column.

```
CREATE VIEW view_d_songs AS
SELECT ID, title AS "Song Title", artist
FROM DJs_on_Demand
WHERE type_code = 'New Age';
```

3. SELECT * FROM view_d_songs. What was returned?

When executed, this query returns the ID, title (aliased as "Song Title"), and artist columns from the DJs on Demand table for each record where the type_code is 'New Age'.

4. REPLACE view_d_songs. Add type_code to the column list. Use aliases for all columns.

```
CREATE OR REPLACE VIEW view_d_songs AS
SELECT ID AS "Song ID", title AS "Song Title", artist, type_code
FROM DJs_on_Demand
WHERE type_code = 'New Age';
```

Or use alias after the CREATE statement as shown.

5. Jason Tsang, the disk jockey for DJs on Demand, needs a list of the past events and those planned for the coming months so he can make arrangements for each event's equipment setup. As the company manager, you do not want him to have access to the price that clients paid for their events. Create a view for Jason to use that displays the name of the event, the event date, and the theme description. Use aliases for each column name.

```
CREATE VIEW event_list AS
SELECT event_name AS "Event Name", event_date AS "Event Date", theme_description AS "Theme
Description"
FROM events;
```

6. It is company policy that only upper-level management be allowed access to individual employee salaries. The department managers, however, need to know the minimum, maximum, and average salaries, grouped by department. Use the Oracle database to prepare a view that displays the needed information for department managers.

```
CREATE VIEW department_salary_summary AS
SELECT department_id,
       MIN(salary) AS min_salary,
       MAX(salary) AS max_salary,
       AVG(salary) AS avg_salary
FROM employees
GROUP BY department_id;
```

DML Operations and Views

Use the DESCRIBE statement to verify that you have tables named copy_d_songs, copy_d_events, copy_d_cds, and copy_d_clients in your schema. If you don't, write a query to create a copy of each.

1. Query the data dictionary USER_UPDATABLE_COLUMNS to make sure the columns in the base tables will allow UPDATE, INSERT, or DELETE. All table names in the data dictionary are stored in uppercase.

```
SELECT table_name FROM user_tables WHERE table_name IN ('COPY_D_SONGS', 'COPY_D_EVENTS',  
'COPY_D_CDS', 'COPY_D_CLIENTS');
```

```
SELECT table_name, column_name, updatable FROM user_updatable_columns WHERE table_name IN  
('COPY_D_SONGS', 'COPY_D_EVENTS', 'COPY_D_CDS', 'COPY_D_CLIENTS');
```

Use the same syntax but change table_name of the other tables.

2. Use the CREATE or REPLACE option to create a view of *all* the columns in the copy_d_songs table called view_copy_d_songs.

```
CREATE OR REPLACE VIEW view_copy_d_songs AS  
SELECT * FROM copy_d_songs;
```

3. Use view_copy_d_songs to INSERT the following data into the underlying copy_d_songs table. Execute a SELECT * from copy_d_songs to verify your DML command. See the graphic.

ID	TITLE	DURATION	ARTIST	TYPE_CODE
88	Mello Jello	2	The What	4

```
INSERT INTO view_copy_d_songs (ID, TITLE, DURATION, ARTIST, TYPE_CODE)
VALUES (88, 'Mello Jello', 2, 'The What', 4);
```

-- Verify insertion

```
SELECT * FROM copy_d_songs;
```

4. Create a view based on the DJs on Demand COPY_D_CDS table. Name the view read_copy_d_cds. Select all columns to be included in the view. Add a WHERE clause to restrict the year to 2000. Add the WITH READ ONLY option.

```
CREATE OR REPLACE VIEW read_copy_d_cds AS
SELECT * FROM copy_d_cds
WHERE year = 2000
WITH READ ONLY;
```

5. Using the read_copy_d_cds view, execute a DELETE FROM read_copy_d_cds WHERE cd_number = 90;

```
DELETE FROM read_copy_d_cds WHERE cd_number = 90;
```

6. Use REPLACE to modify read_copy_d_cds. Replace the READ ONLY option with WITH CHECK OPTION CONSTRAINT ck_read_copy_d_cds. Execute a SELECT * statement to verify that the view exists.

```
CREATE OR REPLACE VIEW read_copy_d_cds AS
SELECT * FROM copy_d_cds
WHERE year = 2000
WITH CHECK OPTION CONSTRAINT ck_read_copy_d_cds;
```

7. Use the read_copy_d_cds view to delete any CD of year 2000 from the underlying copy_d_cds.

```
DELETE FROM read_copy_d_cds;
```

8. Use the read_copy_d_cds view to delete cd_number 90 from the underlying copy_d_cds table.

```
DELETE FROM read_copy_d_cds WHERE cd_number = 90;
```

9. Use the read_copy_d_cds view to delete year 2001 records.

```
DELETE FROM read_copy_d_cds WHERE year = 2001;
```

10. Execute a SELECT * statement for the base table copy_d_cds. What rows were deleted?

```
SELECT * FROM copy_d_cds;
```

11. What are the restrictions on modifying data through a view?

Views may not cover all columns of a base table.

Views with certain functions, expressions, or groupings might not be updatable.

Views based on multiple tables or with joins might not be updatable.

Views with set operations like UNION might not be updatable.

12. What is Moore's Law? Do you consider that it will continue to apply indefinitely? Support your opinion with research from the internet.

Moore's Law is an observation made by Gordon Moore, co-founder of Intel, stating that the number of transistors on integrated circuits doubles approximately every two years.

Some argue that physical limitations such as the size of atoms and heat dissipation will eventually halt the trend. Others believe that innovations in materials, architecture, and computing paradigms may extend its validity.

13. What is the "singularity" in terms of computing?

The "Singularity" in computing refers to a hypothetical future event where technological growth becomes uncontrollable and irreversible. It is often associated with the emergence of superintelligent AI or the merger of human intelligence with artificial intelligence.

Managing Views

1. Create a view from the copy_d_songs table called view_copy_d_songs that includes only the title and artist. Execute a SELECT * statement to verify that the view exists.

```
CREATE VIEW view_copy_d_songs AS  
SELECT title, artist FROM copy_d_songs;
```

```
SELECT * FROM view_copy_d_songs;
```

2. Issue a DROP view_copy_d_songs. Execute a SELECT * statement to verify that the view has been deleted.

```
DROP VIEW view_copy_d_songs;
```

```
SELECT * FROM view_copy_d_songs;
```

3. Create a query that selects the last name and salary from the Oracle database. Rank the salaries from highest to lowest for the top three employees.

```
SELECT last_name, salary  
FROM (  
    SELECT last_name, salary, RANK() OVER (ORDER BY salary DESC) AS salary_rank  
    FROM employees  
)  
WHERE salary_rank <= 3;
```

4. Construct an inline view from the Oracle database that lists the last name, salary, department ID, and maximum salary for each department. Hint: One query will need to calculate maximum salary by department ID.

```
SELECT last_name, salary, department_id, max_salary  
FROM (  
    SELECT last_name, salary, department_id,  
        MAX(salary) OVER (PARTITION BY department_id) AS max_salary  
    FROM employees  
);
```

5. Create a query that will return the staff members of Global Fast Foods ranked by salary from lowest to highest.

```
SELECT last_name, salary  
FROM global_staff  
ORDER BY salary;
```

Indexes and Synonyms

1. What is an index and what is it used for?

An index in a database is a data structure that improves the speed of data retrieval operations on a database table at the cost of additional writes and storage space.

It is used to quickly locate and access specific rows of a table based on the values of one or more columns.

2. What is a ROWID, and how is it used?

ROWID is a pseudo-column in Oracle that uniquely identifies a row in a database table.

It is used internally by Oracle to quickly locate rows and retrieve their data.

It is not recommended for application use as it is subject to change during certain operations like export/import, partitioning, etc.

3. When will an index be created automatically?

Oracle will automatically create indexes in the following scenarios:

When a PRIMARY KEY or UNIQUE constraint is defined on a table.

When a FOREIGN KEY constraint is defined on a table.

When a column is declared as PRIMARY KEY or UNIQUE.

When an indexed column is used in a query's WHERE clause

4. Create a nonunique index (foreign key) for the DJs on Demand column (cd_number) in the D_TRACK_LISTINGS table. Use the Oracle Application Express SQL Workshop Data Browser to confirm that the index was created.

```
CREATE INDEX idx_cd_number ON d_track_listings(cd_number);
```

5. Use the join statement to display the indexes and uniqueness that exist in the data dictionary for the DJs on Demand D_SONGS table.

```
SELECT index_name, uniqueness  
FROM user_indexes  
WHERE table_name = 'D_SONGS';
```

6. Use a SELECT statement to display the index_name, table_name, and uniqueness from the data dictionary USER_INDEXES for the DJs on Demand D_EVENTS table.

```
SELECT index_name, table_name, uniqueness  
FROM user_indexes  
WHERE table_name = 'D_EVENTS';
```

7. Write a query to create a synonym called dj_tracks for the DJs on Demand d_track_listings table.

```
CREATE SYNONYM dj_tracks FOR d_track_listings;
```


8. Create a function-based index for the last_name column in DJs on Demand D_PARTNERS table that makes it possible not to have to capitalize the table name for searches. Write a SELECT statement that would use this index.

```
CREATE INDEX idx_last_name_lowercase ON d_partners(LOWER(last_name));
```

```
SELECT *  
FROM d_partners  
WHERE LOWER(last_name) = 'smith';
```

9. Create a synonym for the D_TRACK_LISTINGS table. Confirm that it has been created by querying the data dictionary.

```
CREATE SYNONYM track_listings_syn FOR
```

```
d_track_listings;
```

```
SELECT * FROM all_synonyms WHERE synonym_name  
= 'TRACK_LISTINGS_SYN';
```

10. Drop the synonym that you created in question

```
DROP SYNONYM track_listings_syn;
```

Evaluation Procedure	Marks awarded
Query(5)	
Execution (5)	
Viva(5)	
Total (15)	
Faculty Signature	

EXERCISE-14

OTHER DATABASE OBJECTS

Objectives

After the completion of this exercise, the students will be able to do the following:

- Create, maintain, and use sequences
- Create and maintain indexes

Database Objects

Many applications require the use of unique numbers as primary key values. You can either build code into the application to handle this requirement or use a sequence to generate unique numbers. If you want to improve the performance of some queries, you should consider creating an index. You

can also use indexes to enforce uniqueness on a column or a collection of columns. You can provide alternative names for objects by using synonyms.

What Is a Sequence?

A sequence:

- Automatically generates unique numbers
- Is a sharable object
- Is typically used to create a primary key value
- Replaces application code
- Speeds up the efficiency of accessing sequence values when cached in memory

The CREATE SEQUENCE Statement Syntax

Define a sequence to generate sequential numbers automatically:

```
CREATE SEQUENCE sequence
[INCREMENT BY n]
[START WITH n]
[{MAXVALUE n | NOMAXVALUE}]
[{MINVALUE n | NOMINVALUE}]
[{CYCLE | NOCYCLE}]
[{CACHE n | NOCACHE}];
```

In the syntax:

sequence is the name of the sequence generator

INCREMENT BY *n* specifies the interval between sequence numbers where *n* is an integer (If this clause is omitted, the sequence increments by 1.)

START WITH *n* specifies the first sequence number to be generated (If this clause is omitted, the sequence starts with 1.)

MAXVALUE *n* specifies the maximum value the sequence can generate

NOMAXVALUE specifies a maximum value of 10^{27} for an ascending sequence and -1 for a descending sequence (This is the default option.)

MINVALUE *n* specifies the minimum sequence value

NOMINVALUE specifies a minimum value of 1 for an ascending sequence and $-(10^{26})$ for a descending sequence (This is the default option.)

CYCLE | NOCYCLE specifies whether the sequence continues to generate values after reaching its maximum or minimum value (NOCYCLE is the default option.)

CACHE *n* | NOCACHE specifies how many values the Oracle server preallocates and keep in memory (By default, the Oracle server caches 20 values.)

Creating a Sequence

- Create a sequence named DEPT_DEPTID_SEQ to be used for the primary key of the DEPARTMENTS table.
- Do not use the CYCLE option.

EXAMPLE:

```
CREATE SEQUENCE dept_deptid_seq  
INCREMENT BY 10  
START WITH 120  
MAXVALUE 9999  
NOCACHE  
NOCYCLE;
```

Confirming Sequences

- Verify your sequence values in the USER_SEQUENCES data dictionary table.
- The LAST_NUMBER column displays the next available sequence number if NOCACHE is specified.

EXAMPLE:

```
SELECT sequence_name, min_value, max_value, increment_by, last_number
```

NEXTVAL and CURRVAL Pseudocolumns

- NEXTVAL returns the next available sequence value. It returns a unique value every time it is referenced, even for different users.
- CURRVAL obtains the current sequence value.
- NEXTVAL must be issued for that sequence before CURRVAL contains a value.

Rules for Using NEXTVAL and CURRVAL

You can use NEXTVAL and CURRVAL in the following contexts:

- The SELECT list of a SELECT statement that is not part of a subquery
- The SELECT list of a subquery in an INSERT statement
- The VALUES clause of an INSERT statement
- The SET clause of an UPDATE statement

You cannot use NEXTVAL and CURRVAL in the following contexts:

- The SELECT list of a view
- A SELECT statement with the DISTINCT keyword

- A SELECT statement with GROUP BY, HAVING, or ORDER BY clauses
- A subquery in a SELECT, DELETE, or UPDATE statement
- The DEFAULT expression in a CREATE TABLE or ALTER TABLE statement

Using a Sequence

- Insert a new department named “Support” in location ID 2500.
- View the current value for the DEPT_DEPTID_SEQ sequence.

EXAMPLE:

```
INSERT INTO departments(department_id, department_name, location_id)
VALUES (dept_deptid_seq.NEXTVAL, 'Support', 2500);
```

```
SELECT dept_deptid_seq.CURRVAL FROM dual;
```

The example inserts a new department in the DEPARTMENTS table. It uses the DEPT_DEPTID_SEQ sequence for generating a new department number as follows:

You can view the current value of the sequence:

```
SELECT dept_deptid_seq.CURRVAL FROM dual;
```

Removing a Sequence

- Remove a sequence from the data dictionary by using the DROP SEQUENCE statement.
- Once removed, the sequence can no longer be referenced.

EXAMPLE:

```
DROP SEQUENCE dept_deptid_seq;
```

What is an Index?

An index:

- Is a schema object
- Is used by the Oracle server to speed up the retrieval of rows by using a pointer
- Can reduce disk I/O by using a rapid path access method to locate data quickly
- Is independent of the table it indexes
- Is used and maintained automatically by the Oracle server

How Are Indexes Created?

- Automatically: A unique index is created automatically when you define a PRIMARY KEY or UNIQUE constraint in a table definition.
- Manually: Users can create nonunique indexes on columns to speed up access to the rows.

Types of Indexes

Two types of indexes can be created. One type is a unique index: the Oracle server automatically creates this index when you define a column in a table to have a PRIMARY KEY or a UNIQUE key constraint. The name of the index is the name given to the constraint.

The other type of index is a nonunique index, which a user can create. For example, you can create a FOREIGN KEY column index for a join in a query to improve retrieval speed.

Creating an Index

- Create an index on one or more columns.
- Improve the speed of query access to the LAST_NAME column in the EMPLOYEES table.

```
CREATE INDEX index
ON table (column[, column]...);
```

EXAMPLE:

```
CREATE INDEX emp_last_name_idx
ON employees(last_name);
```

In the syntax:

index is the name of the index

table is the name of the table

column is the name of the column in the table to be indexed

When to Create an Index

You should create an index if:

- A column contains a wide range of values
- A column contains a large number of null values
- One or more columns are frequently used together in a WHERE clause or a join condition
- The table is large and most queries are expected to retrieve less than 2 to 4 percent of the rows

When Not to Create an Index

It is usually not worth creating an index if:

- The table is small
- The columns are not often used as a condition in the query
- Most queries are expected to retrieve more than 2 to 4 percent of the rows in the table
- The table is updated frequently
- The indexed columns are referenced as part of an Expression

Confirming Indexes

- The USER_INDEXES data dictionary view contains the name of the index and its uniqueness.
- The USER_IND_COLUMNS view contains the index name, the table name, and the column name.

EXAMPLE:

```
SELECT ic.index_name, ic.column_name, ic.column_position col_pos, ix.uniqueness
FROM user_indexes ix, user_ind_columns ic
WHERE ic.index_name = ix.index_name
AND ic.table_name = 'EMPLOYEES';
```

Removing an Index

- Remove an index from the data dictionary by using the DROP INDEX command.
- Remove the UPPER_LAST_NAME_IDX index from the data dictionary.
- To drop an index, you must be the owner of the index or have the DROP ANY INDEX privilege.

```
DROP INDEX upper_last_name_idx;
```

```
DROP INDEX index;
```

Find the Solution for the following:

1. Create a sequence to be used with the primary key column of the DEPT table. The sequence should start at 200 and have a maximum value of 1000. Have your sequence increment by ten numbers. Name the sequence DEPT_ID_SEQ.

```
CREATE SEQUENCE DEPT_ID_SEQ
START WITH 200
INCREMENT BY 10
MAXVALUE 1000
NOCYCLE;
```

2. Write a query in a script to display the following information about your sequences: sequence name, maximum value, increment size, and last number

```
SELECT sequence_name, max_value, increment_by, last_number
FROM user_sequences;
```

3. Write a script to insert two rows into the DEPT table. Name your script lab12_3.sql. Be sure to use the sequence that you created for the ID column. Add two departments named Education and Administration. Confirm your additions. Run the commands in your script.

```
-- lab12_3.sql
INSERT INTO DEPT (DEPT_ID, DEPT_NAME)
VALUES (DEPT_ID_SEQ.NEXTVAL, 'Education');

INSERT INTO DEPT (DEPT_ID, DEPT_NAME)
VALUES (DEPT_ID_SEQ.NEXTVAL, 'Administration');

SELECT * FROM DEPT;

@lab12_3.sql
```

4. Create a nonunique index on the foreign key column (DEPT_ID) in the EMP table.

```
CREATE INDEX idx_dept_id ON EMP(DEPT_ID);
```

5. Display the indexes and uniqueness that exist in the data dictionary for the EMP table.

```
SELECT index_name, uniqueness
FROM user_indexes
WHERE table_name = 'EMP';
```

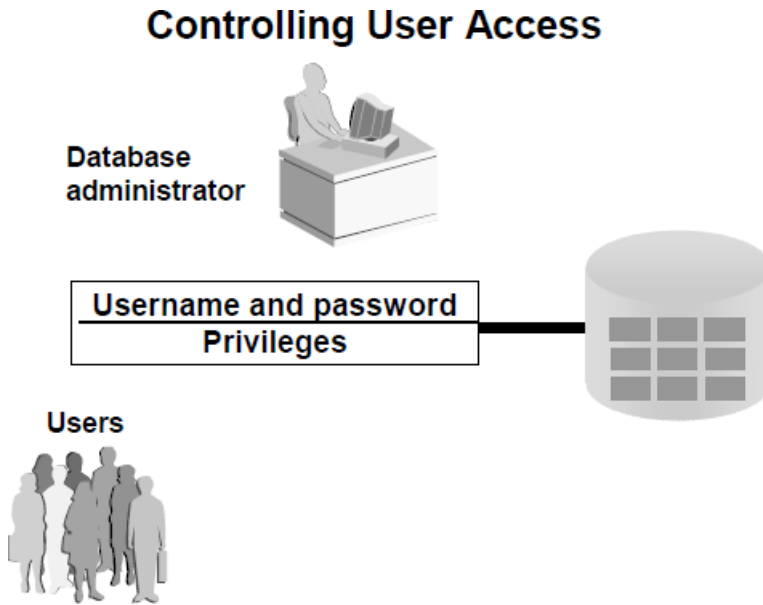
EXERCISE-15

Controlling User Access

Objectives

After the completion of this exercise, the students will be able to do the following:

- Create users
- Create roles to ease setup and maintenance of the security model
- Use the GRANT and REVOKE statements to grant and revoke object privileges
- Create and access database links



Controlling User Access

In a multiple-user environment, you want to maintain security of the database access and use. With Oracle server database security, you can do the following:

- Control database access
- Give access to specific objects in the database
- Confirm given and received *privileges* with the Oracle data dictionary
- Create synonyms for database objects

Privileges

- Database security:
 - System security
 - Data security
- System privileges: Gaining access to the database
- Object privileges: Manipulating the content of the database objects
- Schemas: Collections of objects, such as tables, views, and sequences

System Privileges

- More than 100 privileges are available.
- The database administrator has high-level system privileges for tasks such as:
 - Creating new users

- Removing users
- Removing tables
- Backing up tables

Typical DBA Privileges

System Privilege	Operations Authorized
CREATE USER	Grantee can create other Oracle users (a privilege required for a DBA role).
DROP USER	Grantee can drop another user.
DROP ANY TABLE	Grantee can drop a table in any schema.
BACKUP ANY TABLE	Grantee can back up any table in any schema with the export utility.
SELECT ANY TABLE	Grantee can query tables, views, or snapshots in any schema.
CREATE ANY TABLE	Grantee can create tables in any schema.

Creating Users

The DBA creates users by using the CREATE USER statement.

EXAMPLE:

```
CREATE USER scott IDENTIFIED BY tiger;
```

User System Privileges

- Once a user is created, the DBA can grant specific system privileges to a user.
- An application developer, for example, may have the following system privileges:

- CREATE SESSION
- CREATE TABLE
- CREATE SEQUENCE
- CREATE VIEW
- CREATE PROCEDURE

GRANT *privilege* [, *privilege*...]

TO *user* [, *user*/ *role*, *PUBLIC*...];

Typical User Privileges

System Privilege	Operations Authorized
CREATE SESSION	Connect to the database
CREATE TABLE	Create tables in the user's schema
CREATE SEQUENCE	Create a sequence in the user's schema
CREATE VIEW	Create a view in the user's schema
CREATE PROCEDURE	Create a stored procedure, function, or package in the user's schema

In the syntax:

privilege is the system privilege to be granted

user |*role*|*PUBLIC* is the name of the user, the name of the role, or *PUBLIC* designates that every user is granted the privilege

Note: Current system privileges can be found in the dictionary view SESSION_PRIVS.

Granting System Privileges

The DBA can grant a user specific system privileges.

GRANT create session, create table, create sequence, create view TO scott;

What is a Role?

A role is a named group of related privileges that can be granted to the user. This method makes it easier to revoke and maintain privileges.

A user can have access to several roles, and several users can be assigned the same role. Roles are typically created for a database application.

Creating and Assigning a Role

First, the DBA must create the role. Then the DBA can assign privileges to the role and users to the role.

Syntax

CREATE ROLE *role*;

In the syntax:

role is the name of the role to be created

Now that the role is created, the DBA can use the GRANT statement to assign users to the role as well as assign privileges to the role.

Creating and Granting Privileges to a Role

CREATE ROLE manager;

Role created.

GRANT create table, create view TO manager;

Grant succeeded.

GRANT manager TO DEHAAN, KOCHHAR;

Grant succeeded.

- Create a role
- Grant privileges to a role
- Grant a role to users

Changing Your Password

- The DBA creates your user account and initializes your password.

- You can change your password by using the

ALTER USER statement.

ALTER USER scott

IDENTIFIED BY lion;

User altered.

Object Privileges

Object Privilege	Table	View	Sequence	Procedure
ALTER	√		√	
DELETE	√	√		
EXECUTE				√
INDEX	√			
INSERT	√	√		
REFERENCES	√	√		
SELECT	√	√	√	
UPDATE	√	√		

Object Privileges

- Object privileges vary from object to object.
- An owner has all the privileges on the object.
- An owner can give specific privileges on that owner's object.

GRANT *object_priv* [(*columns*)]

ON *object*

TO {*user*|*role*|PUBLIC}

[WITH GRANT OPTION];

In the syntax:

object_priv is an object privilege to be granted

ALL specifies all object privileges

columns specifies the column from a table or view on which privileges are granted

ON *object* is the object on which the privileges are granted

TO identifies to whom the privilege is granted

PUBLIC grants object privileges to all users

WITH GRANT OPTION allows the grantee to grant the object privileges to other users and roles

Granting Object Privileges

- Grant query privileges on the EMPLOYEES table.
- Grant privileges to update specific columns to users and roles.

GRANT select

ON employees

TO sue, rich;

```
GRANT update (department_name, location_id)
ON departments
TO scott, manager;
```

Using the WITH GRANT OPTION and PUBLIC Keywords

- Give a user authority to pass along privileges.
- Allow all users on the system to query data from Alice's DEPARTMENTS table.

```
GRANT select, insert
ON departments
TO scott
WITH GRANT OPTION;
```

```
.
GRANT select
ON alice.departments
TO PUBLIC;
```

How to Revoke Object Privileges

- You use the REVOKE statement to revoke privileges granted to other users.
 - Privileges granted to others through the WITH GRANT OPTION clause are also revoked.
- ```
REVOKE {privilege [, privilege...]|ALL}
ON object
FROM {user[, user...]|role|PUBLIC}
[CASCADE CONSTRAINTS];
```

#### **In the syntax:**

CASCADE is required to remove any referential integrity constraints made to the CONSTRAINTS object by means of the REFERENCES privilege

### **Revoking Object Privileges**

As user Alice, revoke the SELECT and INSERT privileges given to user Scott on the DEPARTMENTS table.

```
REVOKE select, insert
ON departments
FROM scott;
```

**Find the Solution for the following:**

1. What privilege should a user be given to log on to the Oracle Server? Is this a system or an object privilege?

Users need the "CREATE SESSION" privilege to log onto the Oracle server. It is a system privilege

2. What privilege should a user be given to create tables?

Users need the "CREATE TABLE" privilege to create tables.

3. If you create a table, who can pass along privileges to other users on your table?

The owner of a table can grant privileges to other users on that table.

4. You are the DBA. You are creating many users who require the same system privileges. What should you use to make your job easier?

To make user creation easier for users requiring the same system privileges, you can create roles and grant the necessary privileges to those roles. Then, you can simply assign the roles to the users.

5. What command do you use to change your password?

```
ALTER USER username IDENTIFIED BY new_password;
```

6. Grant another user access to your DEPARTMENTS table. Have the user grant you query access to his or her DEPARTMENTS table.

```
GRANT SELECT ON team1.departments TO your_username;
GRANT SELECT ON your_schema.departments TO team1_username;
```

7. Query all the rows in your DEPARTMENTS table.

```
SELECT * FROM your_schema.departments;
```

8. Add a new row to your DEPARTMENTS table. Team 1 should add Education as department number 500. Team 2 should add Human Resources department number 510. Query the other team's table.

```
INSERT INTO team1.departments (department_id, department_name)
VALUES (500, 'Education');
```

```
INSERT INTO your_schema.departments (department_id, department_name)
VALUES (510, 'Human Resources');
```

```
SELECT * FROM team1.departments; -- For Team 1's table
SELECT * FROM your_schema.departments; -- For Your table
```

9. Query the USER\_TABLES data dictionary to see information about the tables that you own.

```
SELECT * FROM user_tables;
```

10. Revoke the SELECT privilege on your table from the other team.

```
REVOKE SELECT ON your_schema.departments FROM team1_username;
```

11. Remove the row you inserted into the DEPARTMENTS table in step 8 and save the changes.

```
DELETE FROM your_schema.departments WHERE department_id IN (500, 510);
```

| <u>Evaluation Procedure</u>    | <u>Marks awarded</u> |
|--------------------------------|----------------------|
| <u>Practice Evaluation (5)</u> |                      |
| <u>Viva(5)</u>                 |                      |
| <u>Total (10)</u>              |                      |
| <u>Faculty Signature</u>       |                      |

PL/SQL

## **PL/SQL**

### **Control Structures**

In addition to SQL commands, PL/SQL can also process data using flow of statements. The flow of control statements are classified into the following categories.

- Conditional control - Branching
- Iterative control - looping
- Sequential control

#### **BRANCHING in PL/SQL:**

Sequence of statements can be executed on satisfying certain condition.

If statements are being used and different forms of if are:

1. Simple IF
2. ELSIF
3. ELSE IF

#### **SIMPLE IF:**

##### **Syntax:**

IF condition THEN

    statement1;

    statement2;

END IF;

#### **IF-THEN-ELSE STATEMENT:**

##### **Syntax:**

IF condition THEN

    statement1;

ELSE

    statement2;

END IF;

#### **ELSIF STATEMENTS:**

##### **Syntax:**

IF condition1 THEN

    statement1;

ELSIF condition2 THEN

statement2;

ELSIF condition3 THEN

statement3;

ELSE

statementn;

END IF;

**NESTED IF :**

**Syntax:**

IF condition THEN

statement1;

ELSE

IF condition THEN

statement2;

ELSE

statement3;

END IF;

END IF;

ELSE

statement3;

END IF;

**SELECTION IN PL/SQL(Sequential Controls)**

**SIMPLE CASE**

**Syntax:**

CASE SELECTOR

WHEN Expr1 THEN statement1;

WHEN Expr2 THEN statement2;

:



ELSE

Statement n;

END CASE;

### **SEARCHED CASE:**

CASE

WHEN searchcondition1 THEN statement1;

WHEN searchcondition2 THEN statement2;

:

:

ELSE

statementn;

END CASE;

### **ITERATIONS IN PL/SQL**

Sequence of statements can be executed any number of times using loop construct.

It is broadly classified into:

- Simple Loop
- For Loop
- While Loop

### **SIMPLE LOOP**

#### **Syntax:**

LOOP

statement1;

EXIT [ WHEN Condition];

END LOOP;

### **WHILE LOOP**

#### **Syntax:**

WHILE condition LOOP

statement1;

statement2;

END LOOP;

## **FOR LOOP**

### **Syntax:**

FOR counter IN [REVERSE]

    LowerBound..UpperBound

LOOP

statement1;

statement2;

END LOOP;

## PROGRAM 1

Write a PL/SQL block to calculate the incentive of an employee whose ID is 110.

DECLARE

incentive NUMBER(8,2);

BEGIN

SELECT salary \* 0.12 INTO incentive -- Assuming 12% incentive on salary

FROM employees

WHERE employee\_id = 110;

DBMS\_OUTPUT.PUT\_LINE('Incentive = ' || TO\_CHAR(incentive));

END;

/

## PROGRAM 2

Write a PL/SQL block to show an invalid case-insensitive reference to a quoted and without quoted user-defined identifier.

DECLARE

myVar VARCHAR2(20);

BEGIN

-- This will cause an error because "MyVar" is not declared

MyVar := 'This is a value';

DBMS\_OUTPUT.PUT\_LINE(myVar);

END;

/

### PROGRAM 3

Write a PL/SQL block to adjust the salary of the employee whose ID 122.

Sample table: employees

DECLARE

new\_salary NUMBER(8,2);

BEGIN

-- Update salary with a specific value (replace 5000 with desired amount)

UPDATE employees

SET salary = 5000

WHERE employee\_id = 122;

DBMS\_OUTPUT.PUT\_LINE('Salary updated (if employee found).');

END;

/

#### PROGRAM 4

Write a PL/SQL block to create a procedure using the "IS [NOT] NULL Operator" and show AND operator returns TRUE if and only if both operands are TRUE.

DECLARE

var1 VARCHAR2(20) := 'Hello';

var2 VARCHAR2(20);

BEGIN

IF var1 IS NOT NULL AND var2 IS NULL THEN

DBMS\_OUTPUT.PUT\_LINE('var1 has a value, var2 is empty.');

END IF;

END;

/

## PROGRAM 5

Write a PL/SQL block to describe the usage of LIKE operator including wildcard characters and escape character.

DECLARE

emp\_name VARCHAR2(20);

BEGIN

SELECT last\_name

FROM employees

WHERE last\_name LIKE 'Smith%'; -- Matches all names starting with "Smith"

SELECT first\_name

FROM employees

WHERE first\_name LIKE '%en%'; -- Matches names containing "en"

SELECT email

FROM employees

WHERE email LIKE 'user\\_%'; -- Matches emails with "user\_" followed by anything using escape character

DBMS\_OUTPUT.PUT\_LINE('These are examples of LIKE operator usage.');

END;

/

## PROGRAM 6

Write a PL/SQL program to arrange the number of two variable in such a way that the small number will store in num\_small variable and large number will store in num\_large variable.

DECLARE

num1 NUMBER(10);

num2 NUMBER(10);

temp NUMBER(10); -- Temporary variable to swap values

BEGIN

num1 := 20;

num2 := 15;

IF num1 > num2 THEN

temp := num1;

num1 := num2;

num2 := temp;

END IF;

DBMS\_OUTPUT.PUT\_LINE('Smaller number in num\_small: ' || num1);

DBMS\_OUTPUT.PUT\_LINE('Larger number in num\_large: ' || num2);

END;

/



## PROGRAM 7

Write a PL/SQL procedure to calculate the incentive on a target achieved and display the message either the record updated or not.

```
PROCEDURE calculate_incentive (

 p_employee_id IN NUMBER,

 p_target_achieved IN NUMBER,

 p_incentive_rate IN NUMBER -- You can add more parameters if needed

) IS

 v_updated BOOLEAN := FALSE; -- Flag to track update status

 v_incentive NUMBER(8,2);

BEGIN

 -- Update logic with incentive calculation (replace with your specific logic)

 UPDATE employees

 SET commission_earned = commission_earned + p_target_achieved *
p_incentive_rate

 WHERE employee_id = p_employee_id;

 IF SQL%ROWCOUNT > 0 THEN

 v_updated := TRUE;

 END IF;
```

```
v_incentive := p_target_achieved * p_incentive_rate;
```

```
DBMS_OUTPUT.PUT_LINE('Employee ID: ' || p_employee_id);
```

```
IF v_updated THEN
```

```
 DBMS_OUTPUT.PUT_LINE('Record updated. Incentive earned: ' || v_incentive);
```

```
ELSE
```

```
 DBMS_OUTPUT.PUT_LINE('Record not found.');
```

```
END IF;
```

```
END calculate_incentive;
```

```
/
```

```
-- Call the procedure with specific values
```

```
BEGIN
```

```
 calculate_incentive(110, 1000, 0.1); -- Replace with your data
```

```
END;
```

```
/
```

## PROGRAM 8

Write a PL/SQL procedure to calculate incentive achieved according to the specific sale limit.

```
PROCEDURE calculate_incentive (
```

```
 p_sales_amount IN NUMBER
```

```
) IS
```

```
 v_incentive NUMBER(8,2) := 0; -- Initialize incentive to 0
```

```
BEGIN
```

```
 -- Define incentive rates and limits (replace with your actual data)
```

```
 incentive_rates := RECORD (
```

```
 sale_limit NUMBER,
```

```
 rate NUMBER
```

```
);
```

```
 incentive_rates.sale_limit(1) := 1000;
```

```
 incentive_rates.rate(1) := 0.1; -- 10% incentive for sales up to 1000
```

```
 incentive_rates.sale_limit(2) := 5000;
```

```
 incentive_rates.rate(2) := 0.15; -- 15% incentive for sales between 1000 and 5000
```

```
incentive_rates.sale_limit(3) := 10000; -- You can add more tiers as needed
```

```
incentive_rates.rate(3) := 0.2; -- 20% incentive for sales above 10000
```

```
-- Find the applicable incentive rate based on sales amount
```

```
FOR i IN 1..incentive_rates.COUNT LOOP
```

```
 IF p_sales_amount <= incentive_rates.sale_limit(i) THEN
```

```
 v_incentive := p_sales_amount * incentive_rates.rate(i);
```

```
 EXIT; -- Exit the loop once a matching limit is found
```

```
 END IF;
```

```
END LOOP;
```

```
DBMS_OUTPUT.PUT_LINE('Sales amount: ' || p_sales_amount);
```

```
DBMS_OUTPUT.PUT_LINE('Incentive earned: ' || v_incentive);
```

```
END calculate_incentive;
```

```
/
```

```
-- Call the procedure with a sample sales amount
```

```
BEGIN
```

```
 calculate_incentive(2000);
```

```
END;
```

```
/
```

## PROGRAM 9

Write a PL/SQL program to count number of employees in department 50 and check whether this department have any vacancies or not. There are 45 vacancies in this department.

DECLARE

total\_employees NUMBER;

vacancies CONSTANT NUMBER := 45; -- Replace with the actual number of vacancies

BEGIN

SELECT COUNT(\*) INTO total\_employees

FROM employees

WHERE department\_id = 50;

DBMS\_OUTPUT.PUT\_LINE('The employees are in the department 50: ' || total\_employees);

IF total\_employees >= vacancies THEN

DBMS\_OUTPUT.PUT\_LINE('There are no vacancies in the department 50.');

ELSE

DBMS\_OUTPUT.PUT\_LINE('There are ' || vacancies - total\_employees || ' vacancies in department 50.');

END IF;

END; /

## PROGRAM 10

Write a PL/SQL program to count number of employees in a specific department and check whether this department have any vacancies or not. If any vacancies, how many vacancies are in that department.

DECLARE

department\_id NUMBER, total\_employees NUMBER;

vacancies CONSTANT NUMBER := 45; -- Replace with the actual number of vacancies

BEGIN

DBMS\_OUTPUT.PUT\_LINE('Enter department ID: ');

department\_id := TO\_NUMBER(DBMS\_INPUT.GET\_LINE);

SELECT COUNT(\*) INTO total\_employees FROM employees WHERE  
department\_id = department\_id;

DBMS\_OUTPUT.PUT\_LINE('The employees are in the department ' ||  
department\_id || ': ' || total\_employees);

IF total\_employees >= vacancies THEN

DBMS\_OUTPUT.PUT\_LINE('There are no vacancies in the department ' ||  
department\_id || '.');

ELSE

DBMS\_OUTPUT.PUT\_LINE('There are ' || vacancies - total\_employees || '  
vacancies in department ' || department\_id || '.');

END IF;

END;

## PROGRAM 11

Write a PL/SQL program to display the employee IDs, names, job titles, hire dates, and salaries of all employees.

DECLARE

CURSOR emp\_cursor IS

SELECT e.employee\_id,

e.first\_name || ' ' || e.last\_name AS full\_name,

j.job\_title,

e.hire\_date,

e.salary FROM employees e

INNER JOIN jobs j ON e.job\_id = j.job\_id;

emp\_record emp\_cursor%ROWTYPE;

BEGIN

OPEN emp\_cursor;

LOOP

FETCH emp\_cursor INTO emp\_record;

EXIT WHEN emp\_cursor%NOTFOUND;

DBMS\_OUTPUT.PUT\_LINE('Employee ID: ' || emp\_record.employee\_id);

DBMS\_OUTPUT.PUT\_LINE('Name: ' || emp\_record.full\_name);

DBMS\_OUTPUT.PUT\_LINE('Job Title: ' || emp\_record.job\_title);

```
DBMS_OUTPUT.PUT_LINE('Hire Date: ' || TO_CHAR(emp_record.hire_date,
'DD/MM/YYYY')); -- Format hire date
```

```
DBMS_OUTPUT.PUT_LINE('Salary: ' || emp_record.salary);
```

```
DBMS_OUTPUT.PUT_LINE(' ----- ');
```

```
END LOOP;
```

```
CLOSE emp_cursor;
```

```
END;
```

```
/
```



## PROGRAM 12

Write a PL/SQL program to display the employee IDs, names, and department names of all employees.

DECLARE

CURSOR emp\_cursor IS SELECT e.employee\_id,

e.first\_name || ' ' || e.last\_name AS full\_name,

d.department\_name FROM employees e INNER JOIN departments d ON  
e.department\_id = d.department\_id;

emp\_record emp\_cursor%ROWTYPE;

BEGIN

OPEN emp\_cursor;

LOOP

FETCH emp\_cursor INTO emp\_record;

EXIT WHEN emp\_cursor%NOTFOUND;

DBMS\_OUTPUT.PUT\_LINE('Employee ID: ' || emp\_record.employee\_id);

DBMS\_OUTPUT.PUT\_LINE('Name: ' || emp\_record.full\_name);

DBMS\_OUTPUT.PUT\_LINE('Department: ' || emp\_record.department\_name);

DBMS\_OUTPUT.PUT\_LINE(' ----- ');

END LOOP;

CLOSE emp\_cursor;

END;

## PROGRAM 13

Write a PL/SQL program to display the job IDs, titles, and minimum salaries of all jobs.

DECLARE

CURSOR job\_cursor IS

SELECT job\_id, job\_title, min\_salary

FROM jobs;

job\_record job\_cursor%ROWTYPE;

BEGIN

OPEN job\_cursor;

LOOP

FETCH job\_cursor INTO job\_record;

EXIT WHEN job\_cursor%NOTFOUND;

DBMS\_OUTPUT.PUT\_LINE('Job ID: ' || job\_record.job\_id);

DBMS\_OUTPUT.PUT\_LINE('Job Title: ' || job\_record.job\_title);

DBMS\_OUTPUT.PUT\_LINE('Minimum Salary: ' || job\_record.min\_salary);

DBMS\_OUTPUT.PUT\_LINE(' ----- ');

END LOOP;

CLOSE job\_cursor;

END;

/

## PROGRAM 14

Write a PL/SQL program to display the employee IDs, names, and job history start dates of all employees.

DECLARE

CURSOR emp\_history\_cursor IS

```
SELECT e.employee_id, e.first_name || ' ' || e.last_name AS full_name,
 jh.start_date FROM employees e INNER JOIN job_history jh ON
e.employee_id = jh.employee_id;
```

emp\_record emp\_history\_cursor%ROWTYPE;

BEGIN

OPEN emp\_history\_cursor;

LOOP

FETCH emp\_history\_cursor INTO emp\_record;

EXIT WHEN emp\_history\_cursor%NOTFOUND;

DBMS\_OUTPUT.PUT\_LINE('Employee ID: ' || emp\_record.employee\_id);

DBMS\_OUTPUT.PUT\_LINE('Name: ' || emp\_record.full\_name);

DBMS\_OUTPUT.PUT\_LINE('Job History Start Date: ' ||  
TO\_CHAR(emp\_record.start\_date, 'DD/MM/YYYY'));

DBMS\_OUTPUT.PUT\_LINE('-----');

END LOOP;

CLOSE emp\_history\_cursor;

END;

## PROGRAM 15

Write a PL/SQL program to display the employee IDs, names, and job history end dates of all employees.

DECLARE

CURSOR emp\_history\_cursor IS

SELECT e.employee\_id, e.first\_name || ' ' || e.last\_name AS full\_name,

jh.end\_date -- Select end\_date instead of start\_date FROM employees e INNER  
JOIN job\_history jh ON e.employee\_id = jh.employee\_id;

emp\_record emp\_history\_cursor%ROWTYPE;

BEGIN

OPEN emp\_history\_cursor;

LOOP

FETCH emp\_history\_cursor INTO emp\_record;

EXIT WHEN emp\_history\_cursor%NOTFOUND;

DBMS\_OUTPUT.PUT\_LINE('Employee ID: ' || emp\_record.employee\_id);

DBMS\_OUTPUT.PUT\_LINE('Name: ' || emp\_record.full\_name);

DBMS\_OUTPUT.PUT\_LINE('Job History End Date: ' ||  
TO\_CHAR(emp\_record.end\_date, 'DD/MM/YYYY'));

DBMS\_OUTPUT.PUT\_LINE(' ----- ');

END LOOP;CLOSE emp\_history\_cursor;

END;

| <b>Evaluation<br/>Procedure</b>  | <b>Marks awarded</b> |
|----------------------------------|----------------------|
| <b>PL/SQL<br/>Procedure(5)</b>   |                      |
| <b>Program/Execution<br/>(5)</b> |                      |
| <b>Viva(5)</b>                   |                      |
| <b>Total (15)</b>                |                      |
| <b>Faculty Signature</b>         |                      |

## **EXERCISE-16**

### **PROCEDURES AND FUNCTIONS**

#### **PROCEDURES**

##### **DEFINITION**

A procedure or function is a logically grouped set of SQL and PL/SQL statements that perform a specific task. They are essentially sub-programs. Procedures and functions are made up of,

- Declarative part
- Executable part
- Optional exception handling part

These procedures and functions do not show the errors.

##### **KEYWORDS AND THEIR PURPOSES**

**REPLACE:** It recreates the procedure if it already exists.

**PROCEDURE:** It is the name of the procedure to be created.

**ARGUMENT:** It is the name of the argument to the procedure. Paranthesis can be omitted if no arguments are present.

**IN:** Specifies that a value for the argument must be specified when calling the procedure ie. used to pass values to a sub-program. This is the default parameter.

**OUT:** Specifies that the procedure passes a value for this argument back to it's calling environment after execution ie. used to return values to a caller of the sub-program.

**INOUT:** Specifies that a value for the argument must be specified when calling the procedure and that procedure passes a value for this argument back to it's calling environment after execution.

**RETURN:** It is the datatype of the function's return value because every function must return a value, this clause is required.

##### **PROCEDURES – SYNTAX**

```
create or replace procedure <procedure name> (argument {in,out,inout} datatype) {is,as}
variable declaration;
constant declaration;
begin
PL/SQL subprogram body;
exception
exception PL/SQL block;
end;
```

##### **FUNCTIONS – SYNTAX**

```
create or replace function <function name> (argument in datatype,.....) return datatype {is,as}
```

```
variable declaration;
constant declaration;
begin
PL/SQL subprogram body;
exception
exception PL/SQL block;
end;
```

### **CREATING THE TABLE 'ITITEMS' AND DISPLAYING THE CONTENTS**

```
SQL> create table ititems(itemid number(3), actualprice number(5), ordid number(4), prodid
number(4));
Table created.
```

```
SQL> insert into ititems values(101, 2000, 500, 201);
1 row created.
```

```
SQL> insert into ititems values(102, 3000, 1600, 202);
1 row created.
```

```
SQL> insert into ititems values(103, 4000, 600, 202);
1 row created.
```

```
SQL> select * from ititems;
```

| ITEMID | ACTUALPRICE | ORDID | PRODID |
|--------|-------------|-------|--------|
| 101    | 2000        | 500   | 201    |
| 102    | 3000        | 1600  | 202    |
| 103    | 4000        | 600   | 202    |

### **PROGRAM FOR GENERAL PROCEDURE – SELECTED RECORD'S PRICE IS INCREMENTED BY 500 , EXECUTING THE PROCEDURE CREATED AND DISPLAYING THE UPDATED TABLE**

```
SQL> create procedure itsum(identity number, total number) is price number;
2 null_price exception;
3 begin
4 select actualprice into price from ititems where itemid=identity;
5 if price is null then
6 raise null_price;
7 else
8 update ititems set actualprice=actualprice+total where itemid=identity;
9 end if;
10 exception
11 when null_price then
12 dbms_output.put_line('price is null');
13 end;
14 /
Procedure created.
```

```
SQL> exec itsum(101, 500);
PL/SQL procedure successfully completed.
```

```
SQL> select * from ititems;
```

| ITEMID | ACTUALPRICE | ORDID | PRODID |
|--------|-------------|-------|--------|
| 101    | 2500        | 500   | 201    |
| 102    | 3000        | 1600  | 202    |
| 103    | 4000        | 600   | 202    |

### **PROCEDURE FOR 'IN' PARAMETER – CREATION, EXECUTION**

SQL> set serveroutput on;

```
SQL> create procedure yyy (a IN number) is price number;
2 begin
3 select actualprice into price from ititems where itemid=a;
4 dbms_output.put_line('Actual price is ' || price);
5 if price is null then
6 dbms_output.put_line('price is null');
7 end if;
8 end;
9 /
```

Procedure created.

```
SQL> exec yyy(103);
Actual price is 4000
PL/SQL procedure successfully completed.
```

### **PROCEDURE FOR 'OUT' PARAMETER – CREATION, EXECUTION**

SQL> set serveroutput on;

```
SQL> create procedure zzz (a in number, b out number) is identity number;
2 begin
3 select ordid into identity from ititems where itemid=a;
4 if identity<1000 then
5 b:=100;
6 end if;
7 end;
8 /
```

Procedure created.

```
SQL> declare
2 a number;
3 b number;
4 begin
5 zzz(101,b);
6 dbms_output.put_line('The value of b is ' || b);
7 end;
8 /
```

```
The value of b is 100
PL/SQL procedure successfully completed.
```

### **PROCEDURE FOR 'INOUT' PARAMETER – CREATION, EXECUTION**

```
SQL> create procedure itit (a in out number) is
2 begin
```



```
3 a:=a+1;
4 end;
5 /
Procedure created.
```

```
SQL> declare
2 a number:=7;
3 begin
4 itit(a);
5 dbms_output.put_line('The updated value is '||a);
6 end;
7 /
The updated value is 8
PL/SQL procedure successfully completed.
```

### **CREATE THE TABLE 'ITTRAIN' TO BE USED FOR FUNCTIONS**

```
SQL>create table ittrain (tno number(10), tfare number(10));
Table created.
```

```
SQL>insert into ittrain values (1001, 550);
1 row created.
```

```
SQL>insert into ittrain values (1002, 600);
1 row created.
```

```
SQL>select * from ittrain;
 TNO TFARE

1001 550
1002 600
```

### **PROGRAM FOR FUNCTION AND IT'S EXECUTION**

```
SQL> create function aaa (trainnumber number) return number is
2 trainfunction ittrain.tfare % type;
3 begin
4 select tfare into trainfunction from ittrain where tno=trainnumber;
5 return(trainfunction);
6 end;
7 /
```

Function created.

```
SQL> set serveroutput on;
```

```
SQL> declare
2 total number;
3 begin
4 total:=aaa (1001);
5 dbms_output.put_line('Train fare is Rs. '||total);
6 end;
7 /
Train fare is Rs.550
```

PL/SQL procedure successfully completed.

### **Program 1**

#### **FACTORIAL OF A NUMBER USING FUNCTION**

```
CREATE OR REPLACE PROCEDURE calculate_factorial (
 p_number IN NUMBER,
 p_result OUT NUMBER
) IS
BEGIN
 IF p_number < 0 THEN
 DBMS_OUTPUT.PUT_LINE('Error: Factorial is not defined for negative numbers.');
```

    p\_result := NULL; -- Set output parameter to null on error

```
 ELSE
 p_result := 1;
 FOR i IN 1..p_number LOOP
 p_result := p_result * i;
 END LOOP;
 END IF;
END calculate_factorial;
/

-- Example usage: Call the procedure with a number
DECLARE
 factorial_result NUMBER;
BEGIN
 calculate_factorial(5, factorial_result); -- Replace 5 with your desired number
 DBMS_OUTPUT.PUT_LINE('Factorial of 5 is: ' || factorial_result);
END;
/
```

## Program 2

**Write a PL/SQL program using Procedures IN,INOUT,OUT parameters to retrieve the corresponding book information in library**

```
CREATE OR REPLACE PROCEDURE get_book_info (
 p_book_id IN NUMBER,
 p_title OUT VARCHAR2,
 p_author INOUT VARCHAR2 -- Can be used for additional processing like trimming
whitespaces
) IS
 book_record library.books%ROWTYPE;
BEGIN
 -- Query the library.books table
 SELECT title, author
 INTO p_title, book_record.author
 FROM library.books
 WHERE book_id = p_book_id;

 -- Handle cases where book is not found
 IF p_title IS NULL THEN
 DBMS_OUTPUT.PUT_LINE('Book with ID ' || p_book_id || ' not found.');
```

```
 ELSE
 /*
 Optionally process the author name in the INOUT parameter (e.g., trim whitespaces)
 p_author := RTRIM(LTRIM(book_record.author)); -- Example for trimming whitespaces
 */
 END IF;
 END get_book_info;
/

-- Example usage: Call the procedure with a book ID
DECLARE
 book_title VARCHAR2(100);
 author_name VARCHAR2(50);
BEGIN
 author_name := ' F. Scott Fitzgerald '; -- Example with leading/trailing spaces
 get_book_info(123, book_title, author_name); -- Replace 123 with your book ID

 IF book_title IS NOT NULL THEN
 DBMS_OUTPUT.PUT_LINE('Book Title: ' || book_title);
 DBMS_OUTPUT.PUT_LINE('Author: ' || author_name); -- Potentially modified author name
 END IF;
END;
/
```

| <b>Evaluation<br/>Procedure</b>  | <b>Marks awarded</b> |
|----------------------------------|----------------------|
| <b>PL/SQL<br/>Procedure(5)</b>   |                      |
| <b>Program/Execution<br/>(5)</b> |                      |
| <b>Viva(5)</b>                   |                      |
| <b>Total (15)</b>                |                      |
| <b>Faculty Signature</b>         |                      |

## EXERCISE-17

### TRIGGER

#### DEFINITION

A trigger is a statement that is executed automatically by the system as a side effect of a modification to the database. The parts of a trigger are,

- **Trigger statement:** Specifies the DML statements and fires the trigger body. It also specifies the table to which the trigger is associated.
- **Trigger body or trigger action:** It is a PL/SQL block that is executed when the triggering statement is used.
- **Trigger restriction:** Restrictions on the trigger can be achieved

The different uses of triggers are as follows,

- *To generate data automatically*
- *To enforce complex integrity constraints*
- *To customize complex securing authorizations*
- *To maintain the replicate table*
- *To audit data modifications*

#### TYPES OF TRIGGERS

The various types of triggers are as follows,

- **Before:** It fires the trigger before executing the trigger statement.
- **After:** It fires the trigger after executing the trigger statement
- .
- **For each row:** It specifies that the trigger fires once per row
- .
- **For each statement:** This is the default trigger that is invoked. It specifies that the trigger fires once per statement.

#### VARIABLES USED IN TRIGGERS

- :new
- :old

These two variables retain the new and old values of the column updated in the database. The values in these variables can be used in the database triggers for data manipulation

#### SYNTAX

```
create or replace trigger triggername [before/after] {DML statements}
on [tablename] [for each row/statement]
begin
```

-----

-----

-----

### USER DEFINED ERROR MESSAGE

The package “raise\_application\_error” is used to issue the user defined error messages

**Syntax:** raise\_application\_error(error number, 'error message');

The error number can lie between -20000 and -20999.

The error message should be a character string.

## TO CREATE THE TABLE 'ITEMPLS'

```
SQL> create table itempls (ename varchar2(10), eid number(5), salary number(10));
Table created.
```

```
SQL> insert into itempls values('xxx',11,10000);
1 row created.
```

```
SQL> insert into itempls values('yyy',12,10500);
1 row created.
```

```
SQL> insert into itempls values('zzz',13,15500);
1 row created.
```

```
SQL> select * from itempls;
ENAME EID SALARY
```

|     |    |       |
|-----|----|-------|
| xxx | 11 | 10000 |
| yyy | 12 | 10500 |
| zzz | 13 | 15500 |

## TO CREATE A SIMPLE TRIGGER THAT DOES NOT ALLOW INSERT UPDATE AND DELETE OPERATIONS ON THE TABLE

```
SQL> create trigger ittrigg before insert or update or delete on itempls for each row
2 begin
3 raise_application_error(-20010,'You cannot do manipulation');
4 end;
5
6 /
Trigger created.
```

```
SQL> insert into itempls values('aaa',14,34000);
insert into itempls values('aaa',14,34000)
*
```

```
ERROR at line 1:
ORA-20010: You cannot do manipulation
ORA-06512: at "STUDENT.ITTRIGG", line 2
```

ORA-04088: error during execution of trigger 'STUDENT.ITTRIGG'

```
SQL> delete from itempls where ename='xxx';
delete from itempls where ename='xxx'
 *
```

ERROR at line 1:

ORA-20010: You cannot do manipulation

ORA-06512: at "STUDENT.ITTRIGG", line 2

ORA-04088: error during execution of trigger 'STUDENT.ITTRIGG'

```
SQL> update itempls set eid=15 where ename='yyy';
update itempls set eid=15 where ename='yyy'
 *
```

ERROR at line 1:

ORA-20010: You cannot do manipulation

ORA-06512: at "STUDENT.ITTRIGG", line 2

ORA-04088: error during execution of trigger 'STUDENT.ITTRIGG'

### **TO DROP THE CREATED TRIGGER**

```
SQL> drop trigger ittrigg;
```

Trigger dropped.

### **TO CREATE A TRIGGER THAT RAISES AN USER DEFINED ERROR MESSAGE AND DOES NOT ALLOW UPDATION AND INSERTION**

```
SQL> create trigger ittriggs before insert or update of salary on itempls for each row
```

```
2 declare
3 triggsal itempls.salary%type;
4 begin
5 select salary into triggsal from itempls where eid=12;
6 if(:new.salary>triggsal or :new.salary<triggsal) then
7 raise_application_error(-20100,'Salary has not been changed');
8 end if;
9 end;
10 /
```

Trigger created.

```
SQL> insert into itempls values ('bbb',16,45000);
insert into itempls values ('bbb',16,45000)
 *
```

ERROR at line 1:

ORA-04098: trigger 'STUDENT.ITTRIGGS' is invalid and failed re-validation

```
SQL> update itempls set eid=18 where ename='zzz';
update itempls set eid=18 where ename='zzz'
 *
```

ERROR at line 1:

ORA-04298: trigger 'STUDENT.ITTRIGGS' is invalid and failed re-validation

- ☐ Explicit cursor
- ☐ Implicit cursor

### TO CREATE THE TABLE 'SSEMP'

SQL> create table ssemp( eid number(10), ename varchar2(20), job varchar2(20), sal number(10), dnonumber(5));

Table created.

SQL> insert into ssemp values(1,'nala','lecturer',34000,11);

1 row created.

SQL> insert into ssemp values(2,'kala',' seniorlecturer',20000,12);

1 row created.

SQL> insert into ssemp values(5,'ajay','lecturer',30000,11);

1 row created.

SQL> insert into ssemp values(6,'vijay','lecturer',18000,11);

1 row created.

SQL> insert into ssemp values(3,'nila','professor',60000,12);

1 row created.

SQL> select \* from ssemp;

| EID | ENAME | JOB            | SAL   | DNO |
|-----|-------|----------------|-------|-----|
| 1   | nala  | lecturer       | 34000 | 11  |
| 2   | kala  | seniorlecturer | 20000 | 12  |
| 5   | ajay  | lecturer       | 30000 | 11  |
| 6   | vijay | lecturer       | 18000 | 11  |
| 3   | nila  | professor      | 60000 | 12  |

### EXTRA PROGRAMS

#### TO WRITE A PL/SQL BLOCK TO DISPLAY THE EMPLOYEE ID AND EMPLOYEE NAME USING CURSOR FOR LOOP

SQL> set serveroutput on;

SQL> declare

2 begin

3 for emy in (select eid,ename from ssemp)

4 loop

5 dbms\_output.put\_line('Employee id and employee name are '|| emy.eid 'and' || emy.ename);

6 end loop;

7 end;

8 /

Employee id and employee name are 1 and nala

Employee id and employee name are 2 and kala

Employee id and employee name are 5 and ajay

Employee id and employee name are 6 and vijay

Employee id and employee name are 3 and nila

PL/SQL procedure successfully completed.



**TO WRITE A PL/SQL BLOCK TO UPDATE THE SALARY OF ALL EMPLOYEES WHERE DEPARTMENT NO IS 11 BY 5000 USING CURSOR FOR LOOP AND TO DISPLAY THE UPDATED TABLE**

```
SQL> set serveroutput on;
SQL> declare
 2 cursor cem is select eid,ename,sal,dno from ssempp where dno=11;
 3 begin
 4 --open cem;
 5 for rem in cem
 6 loop
 7 update ssempp set sal=rem.sal+5000 where eid=rem.eid;
 8 end loop;
 9 --close cem;
 10 end;
 11 /
```

PL/SQL procedure successfully completed.

```
SQL> select * from ssempp;
```

| EID | ENAME | JOB            | SAL   | DNO |
|-----|-------|----------------|-------|-----|
| 1   | nala  | lecturer       | 39000 | 11  |
| 2   | kala  | seniorlecturer | 20000 | 12  |
| 5   | ajay  | lecturer       | 35000 | 11  |
| 6   | vijay | lecturer       | 23000 | 11  |
| 3   | nila  | professor      | 60000 | 12  |

**TO WRITE A PL/SQL BLOCK TO DISPLAY THE EMPLOYEE ID AND EMPLOYEE NAME WHERE DEPARTMENT NUMBER IS 11 USING EXPLICIT CURSORS**

```
1 declare
2 cursor cenl is select eid,sal from ssempp where dno=11;
3 ecode ssempp.eid%type;
4 esal empp.sal%type;
5 begin
6 open cenl;
7 loop
8 fetch cenl into ecode,esal;
9 exit when cenl%notfound;
10 dbms_output.put_line(' Employee code and employee salary are' || ecode 'and'|| esal);
11 end loop;
12 close cenl;
13* end;
```

```
SQL> /
Employee code and employee salary are 1 and 39000
Employee code and employee salary are 5 and 35000
Employee code and employee salary are 6 and 23000
```

PL/SQL procedure successfully completed.

**TO WRITE A PL/SQL BLOCK TO UPDATE THE SALARY BY 5000 WHERE THE JOB IS LECTURER , TO CHECK IF UPDATES ARE MADE USING IMPLICIT CURSORS AND TO DISPLAY THE UPDATED TABLE**

```
SQL> declare
 2 county number;
 3 begin
 4 update ssemp set sal=sal+10000 where job='lecturer';
 5 county:= sql%rowcount;
 6 if county > 0 then
 7 dbms_output.put_line('The number of rows are '|| county);
 8 end if;
 9 if sql %found then
 10 dbms_output.put_line('Employee record modification successful');
 11 else if sql%notfound then
 12 dbms_output.put_line('Employee record is not found');
 13 end if;
 14 end if;
 15 end;
 16 /
```

The number of rows are 3

Employee record modification successful

PL/SQL procedure successfully completed.

```
SQL> select * from ssemp;
```

| EID | ENAME | JOB            | SAL   | DNO |
|-----|-------|----------------|-------|-----|
| 1   | nala  | lecturer       | 44000 | 11  |
| 2   | kala  | seniorlecturer | 20000 | 12  |
| 5   | ajay  | lecturer       | 40000 | 11  |
| 6   | vijay | lecturer       | 28000 | 11  |
| 3   | nila  | professor      | 60000 | 12  |

**PROGRAMS**

**TO DISPLAY HELLO MESSAGE**

```
SQL> set serveroutput on;
SQL> declare
 2 a varchar2(20);
 3 begin
 4 a:='Hello';
 5 dbms_output.put_line(a);
 6 end;
 7 /
```

Hello

PL/SQL procedure successfully completed.

**TO INPUT A VALUE FROM THE USER AND DISPLAY IT**

```
SQL> set serveroutput on;
SQL> declare
 2 a varchar2(20);
 3 begin
 4 a:=&a;
 5 dbms_output.put_line(a);
 6 end;
 7 /
```

Enter value for a: 5

old 4: a:=&a;

new 4: a:=5;

5

PL/SQL procedure successfully completed.

### **GREATEST OF TWO NUMBERS**

```
SQL> set serveroutput on;
```

```
SQL> declare
 2 a number(7);
 3 b number(7);
 4 begin
 5 a:=&a;
 6 b:=&b;
 7 if(a>b) then
 8 dbms_output.put_line (' The grerater of the two is'|| a);
 9 else
 10 dbms_output.put_line (' The grerater of the two is'|| b);
 11 end if;
 12 end;
 13 /
```

Enter value for a: 5

old 5: a:=&a;

new 5: a:=5;

Enter value for b: 9

old 6: b:=&b;

new 6: b:=9;

The grerater of the two is9

PL/SQL procedure successfully completed.

### **GREATEST OF THREE NUMBERS**

```
SQL> set serveroutput on;
```

```
SQL> declare
 2 a number(7);
 3 b number(7);
 4 c number(7);
 5 begin
 6 a:=&a;
 7 b:=&b;
 8 c:=&c;
```

```
9 if(a>b and a>c) then
10 dbms_output.put_line (' The greatest of the three is ' || a);
11 else if (b>c) then
12 dbms_output.put_line (' The greatest of the three is ' || b);
13 else
14 dbms_output.put_line (' The greatest of the three is ' || c);
15 end if;
16 end if;
17 end;
18 /
```

Enter value for a: 5

old 6: a:=&a;

new 6: a:=5;

Enter value for b: 7

old 7: b:=&b;

new 7: b:=7;

Enter value for c: 1

old 8: c:=&c;

new 8: c:=1;

The greatest of the three is 7

PL/SQL procedure successfully completed.

### **PRINT NUMBERS FROM 1 TO 5 USING SIMPLE LOOP**

SQL> set serveroutput on;

SQL> declare

```
2 a number:=1;
```

```
3 begin
```

```
4 loop
```

```
5 dbms_output.put_line (a);
```

```
6 a:=a+1;
```

```
7 exit when a>5;
```

```
8 end loop;
```

```
9 end;
```

```
10 /
```

1

2

3

4

5

PL/SQL procedure successfully completed.

### **PRINT NUMBERS FROM 1 TO 4 USING WHILE LOOP**

SQL> set serveroutput on;

SQL> declare

```
2 a number:=1;
```

```
3 begin
```

```
4 while(a<5)
5 loop
6 dbms_output.put_line (a);
7 a:=a+1;
8 end loop;
9 end;
10 /
```

1

2

3

4

PL/SQL procedure successfully completed.

### **PRINT NUMBERS FROM 1 TO 5 USING FOR LOOP**

SQL> set serveroutput on;

SQL> declare

```
2 a number:=1;
3 begin
4 for a in 1..5
5 loop
6 dbms_output.put_line (a);
7 end loop;
8 end;
9 /
```

1

2

3

4

5

PL/SQL procedure successfully completed.

### **PRINT NUMBERS FROM 1 TO 5 IN REVERSE ORDER USING FOR LOOP**

SQL> set serveroutput on;

SQL> declare

```
2 a number:=1;
3 begin
4 for a in reverse 1..5
5 loop
6 dbms_output.put_line (a);
7 end loop;
8 end;
9 /
```

5

4

3

2

1

PL/SQL procedure successfully completed.

### **TO CALCULATE AREA OF CIRCLE**

SQL> set serveroutput on;

SQL> declare

```
2 pi constant number(4,2):=3.14;
3 a number(20);
4 r number(20);
5 begin
```

```

6 r:=&r;
7 a:= pi* power(r,2);
8 dbms_output.put_line (' The area of circle is ' || a);
9 end;
10 /

```

Enter value for r: 2

old 6: r:=&r;

new 6: r:=2;

The area of circle is 13

PL/SQL procedure successfully completed.

### **TO CREATE SACCOUNT TABLE**

SQL> create table saccount ( accno number(5), name varchar2(20), bal number(10));

Table created.

SQL> insert into saccount values ( 1,'mala',20000);

1 row created.

SQL> insert into saccount values (2,'kala',30000);

1 row created.

SQL> select \* from saccount;

| ACCNO | NAME | BAL   |
|-------|------|-------|
| 1     | mala | 20000 |
| 2     | kala | 30000 |

SQL> set serveroutput on;

SQL> declare

```

2 a_bal number(7);
3 a_no varchar2(20);
4 debit number(7):=2000;
5 minamt number(7):=500;
6 begin
7 a_no:=&a_no;
8 select bal into a_bal from saccount where accno= a_no;
9 a_bal:= a_bal-debit;
10 if (a_bal > minamt) then
11 update saccount set bal=bal-debit where accno=a_no;
12 end if;
13 end;
14
15 /

```

Enter value for a\_no: 1

old 7: a\_no:=&a\_no;

new 7: a\_no:=1;

PL/SQL procedure successfully completed.

SQL> select \* from saccount;

| ACCNO | NAME | BAL   |
|-------|------|-------|
| 1     | mala | 18000 |
| 2     | kala | 30000 |

### **TO CREATE TABLE SROUTES**

SQL> create table sroutes ( rno number(5), origin varchar2(20), destination varchar2(20), fare

numbe

r(10), distance number(10));

Table created.

```
SQL> insert into sroutes values (2, 'chennai', 'dindugal', 400,230);
1 row created.
SQL> insert into sroutes values (3, 'chennai', 'madurai', 250,300);
1 row created.
SQL> insert into sroutes values (6, 'thanjavur', 'palani', 350,370);
1 row created.
SQL> select * from sroutes;
```

| RNO | ORIGIN    | DESTINATION | FARE | DISTANCE |
|-----|-----------|-------------|------|----------|
| 2   | chennai   | dindugal    | 400  | 230      |
| 3   | chennai   | madurai     | 250  | 300      |
| 6   | thanjavur | palani      | 350  | 370      |

```
SQL> set serveroutput on;
```

```
SQL> declare
2 route sroutes.rno % type;
3 fares sroutes.fare % type;
4 dist sroutes.distance % type;
5 begin
6 route:=&route;
7 select fare, distance into fares , dist from sroutes where rno=route;
8 if (dist < 250) then
9 update sroutes set fare=300 where rno=route;
10 else if dist between 250 and 370 then
11 update sroutes set fare=400 where rno=route;
12 else if (dist > 400) then
13 dbms_output.put_line('Sorry');
14 end if;
15 end if;
16 end if;
17 end;
18 /
Enter value for route: 3
old 6: route:=&route;
new 6: route:=3;
```

PL/SQL procedure successfully completed.

```
SQL> select * from sroutes;
```

| RNO | ORIGIN    | DESTINATION | FARE | DISTANCE |
|-----|-----------|-------------|------|----------|
| 2   | chennai   | dindugal    | 400  | 230      |
| 3   | chennai   | madurai     | 400  | 300      |
| 6   | thanjavur | palani      | 350  | 370      |

### **TO CREATE SCALCULATE TABLE**

```
SQL> create table scalculate (radius number(3), area number(5,2));
Table created.
SQL> desc scalculate;
Name Null? Type
```

RADIUS

AREA

NUMBER(3)

NUMBER(5,2)

SQL> set serveroutput on;

SQL> declare

```
2 pi constant number(4,2):=3.14;
3 area number(5,2);
4 radius number(3);
5 begin
6 radius:=3;
7 while (radius <=7)
8 loop
9 area:= pi* power(radius,2);
10 insert into scalculate values (radius,area);
11 radius:=radius+1;
12 end loop;
13 end;
14 /
```

PL/SQL procedure successfully completed.

SQL> select \* from scalculate;

| RADIUS | AREA |
|--------|------|
|--------|------|

|   |        |
|---|--------|
| 3 | 28.26  |
| 4 | 50.24  |
| 5 | 78.5   |
| 6 | 113.04 |
| 7 | 153.86 |

### **TO CALCULATE FACTORIAL OF A GIVEN NUMBER**

SQL> set serveroutput on;

SQL> declare

```
2 f number(4):=1;
3 i number(4);
4 begin
5 i:=&i;
6 while(i>=1)
7 loop
8 f:=f*i;
9 i:=i-1;
10 end loop;
11 dbms_output.put_line('The value is ' || f);
12 end;
13 /
```

Enter value for i: 5

old 5: i:=&i;

new 5: i:=5;

The value is 120

PL/SQL procedure successfully completed.



### Program 1

Write a code in PL/SQL to develop a trigger that enforces referential integrity by preventing the deletion of a parent record if child records exist.

```
CREATE OR REPLACE TRIGGER prevent_parent_delete
```

```
BEFORE DELETE ON parent_table -- Replace with your parent table name
```

```
FOR EACH ROW
```

```
DECLARE
```

```
 v_child_count NUMBER;
```

```
BEGIN
```

```
 -- Check if any child records exist in the child_table referencing the parent_id
```

```
 SELECT COUNT(*) INTO v_child_count
```

```
 FROM child_table -- Replace with your child table name
```

```
 WHERE child_table.parent_id = :OLD.parent_id;
```

```
 IF v_child_count > 0 THEN
```

```
 RAISE APPLICATION_ERROR(-20001, 'Cannot delete parent record with
existing child records.');
```

```
 END IF;
```

```
END;
```

```
/
```

## Program 2

Write a code in PL/SQL to create a trigger that checks for duplicate values in a specific column and raises an exception if found.

```
CREATE OR REPLACE TRIGGER prevent_duplicate_values

BEFORE INSERT OR UPDATE ON your_table -- Replace with your table name

FOR EACH ROW

DECLARE v_existing_count NUMBER;

BEGIN IF INSERTING THEN

 SELECT COUNT(*) INTO v_existing_count FROM your_table

 WHERE your_column = :NEW.your_column; -- Replace with your column name

ELSE

 SELECT COUNT(*) INTO v_existing_count FROM your_table

 WHERE your_column = :NEW.your_column

 AND your_primary_key <> :NEW.your_primary_key; -- Replace with your
primary key

END IF;

IF v_existing_count > 0 THEN RAISE APPLICATION_ERROR(-20001, 'Duplicate
value found in ' || your_column || '.');

END IF;

END;
```

/

### Program 3

Write a code in PL/SQL to create a trigger that restricts the insertion of new rows if the total of a column's values exceeds a certain threshold.

```
CREATE OR REPLACE TRIGGER limit_column_total
```

```
BEFORE INSERT ON your_table -- Replace with your table name
```

```
FOR EACH ROW
```

```
DECLARE
```

```
 v_current_sum NUMBER;
```

```
 v_threshold NUMBER := 1000; -- Replace with your desired threshold
```

```
BEGIN
```

```
 -- Calculate the current sum of the column excluding the new row
```

```
 SELECT SUM(your_column) INTO v_current_sum -- Replace with your column
name
```

```
 FROM your_table;
```

```
 Check if inserting the new row will exceed the threshold
```

```
 IF v_current_sum + :NEW.your_column > v_threshold THEN
```

```
 RAISE APPLICATION_ERROR(-20001, 'Total value of ' || your_column || '
exceeds the limit.');
```

```
 END IF;
```

```
END;
```

```
/
```

## Program 4

Write a code in PL/SQL to design a trigger that captures changes made to specific columns and logs them in an audit table.

```
CREATE OR REPLACE TRIGGER audit_column_changes
```

```
AFTER UPDATE ON your_table -- Replace with your table name
```

```
FOR EACH ROW
```

```
DECLARE
```

```
 v_old_data your_table%ROWTYPE; -- Replace with your table row type
```

```
 v_audit_record audit_table%ROWTYPE; -- Replace with your audit table row
type
```

```
BEGIN
```

```
 -- Get old row data using pseudo-record for UPDATE
```

```
 SELECT * INTO v_old_data FROM your_table WHERE ROWID = :OLD.ROWID;
```

```
 -- Loop through specific columns to capture changes
```

```
 FOR col IN (SELECT column_name FROM user_tab_cols WHERE table_name =
'YOUR_TABLE' AND audit_flag = 'Y') LOOP -- Replace with your logic for
identifying columns to audit
```

```
 IF v_old_data(col) != :NEW.(col) THEN
```

```
 v_audit_record.table_name := 'YOUR_TABLE'; -- Replace with your table
name
```

```
v_audit_record.column_name := col;
```

```
v_audit_record.old_value := v_old_data(col);
```

```
v_audit_record.new_value := :NEW.(col);
```

```
v_audit_record.audit_date := SYSDATE;
```

```
v_audit_record.user_name := USER;
```

```
-- Insert audit record
```

```
INSERT INTO audit_table VALUES (v_audit_record);
```

```
END IF;
```

```
END LOOP;
```

```
END;
```

```
/
```

## Program 5

Write a code in PL/SQL to implement a trigger that records user activity (inserts, updates, deletes) in an audit log for a given set of tables.

```
CREATE OR REPLACE TRIGGER audit_user_activity
```

```
AFTER INSERT OR UPDATE OR DELETE ON your_table -- Replace with your
first table name
```

```
FOR EACH ROW
```

```
DECLARE
```

```
 v_audit_record audit_log%ROWTYPE; -- Replace with your audit log table row
type
```

```
BEGIN
```

```
 v_audit_record.table_name := SYS_CONTEXT('USERENV', 'TABLE_NAME');
```

```
 v_audit_record.operation := CASE
```

```
 WHEN INSERTING THEN 'INSERT'
```

```
 WHEN UPDATING THEN 'UPDATE'
```

```
 WHEN DELETING THEN 'DELETE'
```

```
 END;
```

```
 v_audit_record.audit_date := SYSDATE;
```

```
 v_audit_record.user_name := USER;
```

```
 -- Replace with logic to identify additional details based on operation (optional)
```

IF INSERTING THEN

-- You can capture specific details for inserts (e.g., :NEW.<column\_name>)

ELSIF UPDATING THEN

-- You can capture specific details for updates (e.g., :OLD.<column\_name>,  
:NEW.<column\_name>)

ELSIF DELETING THEN

-- You can capture specific details for deletes (e.g., :OLD.<column\_name>)

END IF;

-- Insert audit record

INSERT INTO audit\_log VALUES (v\_audit\_record);

END;

/

CREATE OR REPLACE TRIGGER BODY audit\_user\_activity -- Body for  
additional tables (optional)

IS

BEGIN

-- Call the main trigger for other tables you want to audit

audit\_user\_activity;

END;

/

-- Grant execute on the trigger body to tables you want to audit (if using separate body)

GRANT EXECUTE ON audit\_user\_activity TO your\_table2, your\_table3; --

Replace with your table names



## Program 7

Write a code in PL/SQL to implement a trigger that automatically calculates and updates a running total column for a table whenever new rows are inserted.

```
CREATE OR REPLACE TRIGGER update_running_total
```

```
BEFORE INSERT ON your_table -- Replace with your table name
```

```
FOR EACH ROW
```

```
DECLARE
```

```
 v_current_total NUMBER;
```

```
BEGIN
```

```
 -- Check if running_total column is NULL for the new row
```

```
 IF :NEW.running_total IS NULL THEN
```

```
 -- Retrieve the existing maximum running total (if any)
```

```
 SELECT NVL(MAX(running_total), 0) INTO v_current_total
```

```
 FROM your_table;
```

```
 ELSE
```

```
 -- Use the existing value from the new row (assuming updates are handled
elsewhere)
```

```
 v_current_total := :NEW.running_total;
```

```
 END IF;
```

-- Calculate the new running total with the new row's value

v\_current\_total := v\_current\_total + :NEW.your\_value\_column; -- Replace with  
your value column

-- Update the running total in the new row

:NEW.running\_total := v\_current\_total;

END;

/

## Program 8

Write a code in PL/SQL to create a trigger that validates the availability of items before allowing an order to be placed, considering stock levels and pending orders.

```
CREATE OR REPLACE TRIGGER validate_order_availability
BEFORE INSERT ON orders -- Replace with your order table name
FOR EACH ROW
DECLARE
 v_item_available NUMBER := 0;
v_stock_quantity NUMBER;
v_pending_orders NUMBER;
BEGIN
 -- Check if item exists and retrieve stock quantity
 SELECT stock_quantity
 INTO v_stock_quantity
 FROM inventory -- Replace with your inventory table name
 WHERE inventory.item_id = :NEW.item_id;

 IF v_stock_quantity IS NULL THEN
 RAISE APPLICATION_ERROR(-20001, 'Item with ID ' || :NEW.item_id || ' not found in
inventory. ');
 END IF;

 -- Calculate total pending orders (including the new order)
 SELECT SUM(quantity) + :NEW.quantity
 INTO v_pending_orders
 FROM order_items -- Replace with your order items table name
 WHERE order_items.item_id = :NEW.item_id
 AND order_items.order_status IN ('PENDING', 'PROCESSING'); -- Adjust order statuses for
pending orders

 -- Check if sufficient stock is available considering pending orders
 IF v_stock_quantity >= v_pending_orders THEN
 v_item_available := 1;
 END IF;

 IF v_item_available = 0 THEN
 RAISE APPLICATION_ERROR(-20002, 'Insufficient stock or pending orders for item ' ||
:NEW.item_id || '. ');
 END IF;

 -- Insert the order (assuming validation passes)
 INSERT INTO orders (...) VALUES (:NEW...); -- Replace with your actual columns and values
END;
/
```

| <b>Evaluation<br/>Procedure</b>  | <b>Marks awarded</b> |
|----------------------------------|----------------------|
| <b>PL/SQL<br/>Procedure(5)</b>   |                      |
| <b>Program/Execution<br/>(5)</b> |                      |
| <b>Viva(5)</b>                   |                      |
| <b>Total (15)</b>                |                      |
| <b>Faculty Signature</b>         |                      |

# MONGO DB

## MONGO DB

MongoDB is a free and open-source cross-platform document-oriented database. Classified as a NoSQL database, MongoDB avoids the traditional table-based relational database structure in favor of JSON-like documents with dynamic schemas, making the integration of data in certain types of applications easier and faster.

### Create Database using mongosh

After connecting to your database using mongosh, you can see which database you are using by typing db in your terminal.

If you have used the connection string provided from the MongoDB Atlas dashboard, you should be connected to the myFirstDatabase database.

### Show all databases

To see all available databases, in your terminal type show dbs.

Notice that myFirstDatabase is not listed. This is because the database is empty. An empty database is essentially non-existent.

### Change or Create a Database

You can change or create a new database by typing use then the name of the database.

### Create Collection using mongosh

You can create a collection using the createCollection() database method.

### Insert Documents

```
insertOne()
```

```
db.posts.insertOne({
```

```
 title: "Post Title 1",
```

```
 body: "Body of post.",
```

```
 category: "News",
```

```
 likes: 1,
```

```
 tags: ["news", "events"],
```

```
 date: Date()
```

```
})
```

## EXERCISE 18

Structure of 'restaurants' collection:

```
{
 "address": {
 "building": "1007",
 "coord": [-73.856077, 40.848447],
 "street": "Morris Park Ave",
 "zipcode": "10462"
 },
 "borough": "Bronx",
 "cuisine": "Bakery",
 "grades": [
 { "date": { "$date": 1393804800000 }, "grade": "A", "score": 2 },
 { "date": { "$date": 1378857600000 }, "grade": "A", "score": 6 },
 { "date": { "$date": 1358985600000 }, "grade": "A", "score": 10 },
 { "date": { "$date": 1322006400000 }, "grade": "A", "score": 9 },
 { "date": { "$date": 1299715200000 }, "grade": "B", "score": 14 }
],
 "name": "Morris Park Bake Shop",
 "restaurant_id": "30075445"
}
```

**1. Write a MongoDB query to find the restaurant Id, name, borough and cuisine for those restaurants which prepared dish except 'American' and 'Chinees' or restaurant's name begins with letter 'Wil'.**

**db.restaurants.find({**

**\$or: [**

**{ cuisine: { \$nin: ["American", "Chinees"] } },**

**{ name: { \$regex: /^Wil/ } }**

**]**

**}, { restaurant\_id: 1, name: 1, borough: 1, cuisine: 1 })**

2. Write a MongoDB query to find the restaurant Id, name, and grades for those restaurants which achieved a grade of "A" and scored 11 on an ISODate "2014-08-11T00:00:00Z" among many of survey dates..

```
db.restaurants.find({

 "grades.grade": "A",

 "grades.date": ISODate("2014-08-11T00:00:00Z"),

 "grades.score": 11

}, { restaurant_id: 1, name: 1, grades: 1 })
```

3. Write a MongoDB query to find the restaurant Id, name and grades for those restaurants where the 2nd element of grades array contains a grade of "A" and score 9 on an ISODate "2014-08-11T00:00:00Z".

```
db.restaurants.find({

 grades: {

 $elemMatch: {

 grade: "A",

 score: 9,

 date: ISODate("2014-08-11T00:00:00Z")

 }

 }

}, { restaurant_id: 1, name: 1, grades: 1 })
```

**4. Write a MongoDB query to find the restaurant Id, name, address and geographical location for those restaurants where 2nd element of coord array contains a value which is more than 42 and upto 52..**

```
db.restaurants.find({

 "address.coord": { $elemMatch: { $gt: 42, $lt: 52 } }

}, { restaurant_id: 1, name: 1, address: 1 })
```

**5. Write a MongoDB query to arrange the name of the restaurants in ascending order along with all the columns.**

```
db.restaurants.find().sort({ name: 1 })
```

**6. Write a MongoDB query to arrange the name of the restaurants in descending order along with all the columns.**

```
db.restaurants.find().sort({ name: -1 })
```

**7. Write a MongoDB query to arranged the name of the cuisine in ascending order and for that same cuisine borough should be in descending order.**

```
db.restaurants.find().sort({ cuisine: 1, borough: -1 })
```

**8. Write a MongoDB query to know whether all the addresses contains the street or not.**

```
db.restaurants.find({ address: { $not: /street/i } }).count() === 0 // true if all contain
"street"
```

**9. Write a MongoDB query which will select all documents in the restaurants collection where the coord field value is Double.**

```
db.restaurants.find({ "coord": { $type: "double" } })
```

**10. Write a MongoDB query which will select the restaurant Id, name and grades for those restaurants which returns 0 as a remainder after dividing the score by 7.**

```
db.restaurants.find({ "grades.score": { $mod: [7, 0] } })
```



**11. Write a MongoDB query to find the restaurant name, borough, longitude and attitude and cuisine for those restaurants which contains 'mon' as three letters somewhere in its name.**

```
db.restaurants.find({ name: { $regex: /mon/i } }, { name: 1, borough: 1, "address.coord": 1, cuisine: 1 })
```

**12. Write a MongoDB query to find the restaurant name, borough, longitude and latitude and cuisine for those restaurants which contain 'Mad' as first three letters of its name.**

```
db.restaurants.find({ name: { $regex: /^Mad/ } }, { name: 1, borough: 1, "address.coord": 1, cuisine: 1 })
```

**13. Write a MongoDB query to find the restaurants that have at least one grade with a score of less than 5.**

```
db.restaurants.find({ grades: { $elemMatch: { score: { $lt: 5 } } } })
```

**14. Write a MongoDB query to find the restaurants that have at least one grade with a score of less than 5 and that are located in the borough of Manhattan.**

```
db.restaurants.find({

 grades: { $elemMatch: { score: { $lt: 5 } } },

 borough: "Manhattan"

})
```

**15. Write a MongoDB query to find the restaurants that have at least one grade with a score of less than 5 and that are located in the borough of Manhattan or Brooklyn.**

```
db.restaurants.find({

 grades: { $elemMatch: { score: { $lt: 5 } } },

 borough: { $in: ["Manhattan", "Brooklyn"] }

})
```

16. Write a MongoDB query to find the restaurants that have at least one grade with a score of less than 5 and that are located in the borough of Manhattan or Brooklyn, and their cuisine is not American.

```
db.restaurants.find({
```

```
 grades: { $
```

17. Write a MongoDB query to find the restaurants that have at least one grade with a score of less than 5 and that are located in the borough of Manhattan or Brooklyn, and their cuisine is not American or Chinese.

```
db.restaurants.find({
```

```
 $and: [
```

```
 { grades: { $elemMatch: { score: { $lt: 5 } } } },
```

```
 { borough: { $in: ["Manhattan", "Brooklyn"] } },
```

```
 { cuisine: { $nin: ["American", "Chinese"] } }
```

```
]
```

```
}, { restaurant_id: 1, name: 1, borough: 1, cuisine: 1 })
```

18. Write a MongoDB query to find the restaurants that have a grade with a score of 2 and a grade with a score of 6.

```
db.restaurants.find({
```

```
 grades: {
```

```
 $size: { $gte: 2 } // Ensure at least 2 grades
```

```
 },
```

```
 grades: { $elemMatch: { score: { $in: [2, 6] } } } // Grades with scores 2 or 6
```

```
})
```

**19. Write a MongoDB query to find the restaurants that have a grade with a score of 2 and a grade with a score of 6 and are located in the borough of Manhattan.**

```
db.restaurants.find({

 grades: {

 $size: { $gte: 2 } // Ensure at least 2 grades

 },

 grades: { $elemMatch: { score: { $in: [2, 6] } } },

 borough: "Manhattan"

})
```

**20. Write a MongoDB query to find the restaurants that have a grade with a score of 2 and a grade with a score of 6 and are located in the borough of Manhattan or Brooklyn.**

```
db.restaurants.find({

 grades: {

 $size: { $gte: 2 } // Ensure at least 2 grades

 },

 grades: { $elemMatch: { score: { $in: [2, 6] } } },

 borough: { $in: ["Manhattan", "Brooklyn"] }

})
```

**21. Write a MongoDB query to find the restaurants that have a grade with a score of 2 and a grade with a score of 6 and are located in the borough of Manhattan or Brooklyn, and their cuisine is not American.**

```
db.restaurants.find({

 grades: {

 $size: { $gte: 2 } // Ensure at least 2 grades

 },

 grades: { $elemMatch: { score: { $in: [2, 6] } } },

 borough: { $in: ["Manhattan", "Brooklyn"] },

 cuisine: { $ne: "American" }

})
```

**22. Write a MongoDB query to find the restaurants that have a grade with a score of 2 and a grade with a score of 6 and are located in the borough of Manhattan or Brooklyn, and their cuisine is not American or Chinese.**

```
db.restaurants.find({

 grades: {

 $size: { $gte: 2 } // Ensure at least 2 grades

 },

 grades: { $elemMatch: { score: { $in: [2, 6] } } },

 borough: { $in: ["Manhattan", "Brooklyn"] },

 cuisine: { $nin: ["American", "Chinese"] }

})
```

**23. Write a MongoDB query to find the restaurants that have a grade with a score of 2 or a grade with a score of 6.**

```
db.restaurants.find({

 grades: { $elemMatch: { score: { $in: [2, 6] } } }

})
```

### **Sample document of 'movies' collection**

```
{
 _id: ObjectId("573a1390f29313caabcd42e8"),
 plot: 'A group of bandits stage a brazen train hold-up, only to find a determined posse hot on their heels.',
 genres: ['Short', 'Western'],
 runtime: 11,
 cast: [
 'A.C. Abadie',
 "Gilbert M. 'Broncho Billy' Anderson",
 'George Barnes',
 'Justus D. Barnes'
],
 poster: 'https://m.media-amazon.com/images/M/MV5BMTU3NjE5NzYtYTtyNS00MDVmLWlwYjgtMmYwYWlxdDYyNzU2XkEyXkFqcGdeQXVyNzQzNzQxNzI@._V1_SY1000_SX677_AL_.jpg',
 title: 'The Great Train Robbery',
 fullplot: "Among the earliest existing films in American cinema - notable as the first film that presented a narrative story to tell - it depicts a group of cowboy outlaws who hold up a train and rob the passengers. They are then pursued by a Sheriff's posse. Several scenes have color included - all hand tinted.",
 languages: ['English'],
 released: ISODate("1903-12-01T00:00:00.000Z"),
 directors: ['Edwin S. Porter'],
```

```
rated: 'TV-G',
awards: { wins: 1, nominations: 0, text: '1 win.' },
lastupdated: '2015-08-13 00:27:59.177000000',
year: 1903,
imdb: { rating: 7.4, votes: 9847, id: 439 },
countries: ['USA'],
type: 'movie',
tomatoes: {
viewer: { rating: 3.7, numReviews: 2559, meter: 75 },
fresh: 6,
critic: { rating: 7.6, numReviews: 6, meter: 100 },
rotten: 0,
lastUpdated: ISODate("2015-08-08T19:16:10.000Z")
}
```

1. Find all movies with full information from the 'movies' collection that released in the year 1893.

```
db.movies.find({ year: 1893 })
```

2. Find all movies with full information from the 'movies' collection that have a runtime greater than 120 minutes.

```
db.movies.find({ runtime: { $gt: 120 } }) // $gt: greater than
```

3. Find all movies with full information from the 'movies' collection that have "Short" genre.

```
db.movies.find({ genres: "Short" })
```

4. Retrieve all movies from the 'movies' collection that were directed by "William K.L. Dickson" and include complete information for each movie.

```
db.movies.find({ director: "William K.L. Dickson" })
```

5. Retrieve all movies from the 'movies' collection that were released in the USA and include complete information for each movie.

```
db.movies.find({ countries: "USA" })
```

6. Retrieve all movies from the 'movies' collection that have complete information and are rated as "UNRATED".

```
db.movies.find({ rated: "UNRATED" })
```

7. Retrieve all movies from the 'movies' collection that have complete information and have received more than 1000 votes on IMDb.

```
db.movies.find({ votes: { $gt: 1000 } }) // $gt: greater than
```

8. Retrieve all movies from the 'movies' collection that have complete information and have an IMDb rating higher than 7.

```
db.movies.find({ imdb.rating: { $gt: 7 } })
```

9. Retrieve all movies from the 'movies' collection that have complete information and have a viewer rating higher than 4 on Tomatoes.

```
db.movies.find({ "tomatoes.viewerRating": { $gt: 4 } })
```

10. Retrieve all movies from the 'movies' collection that have received an award.

```
db.movies.find({ awards: { $ne: [] } }) // $ne: not empty
```

11. Find all movies with title, languages, released, directors, writers, awards, year, genres, runtime, cast, countries from the 'movies' collection in MongoDB that have at least one nomination.

```
db.movies.find({ awards: { $regex: /Nominated/i } }) // $regex: regular expression, i: case-insensitive
```

12. Find all movies with title, languages, released, directors, writers, awards, year, genres, runtime, cast, countries from the 'movies' collection in MongoDB with cast including "Charles Kayser".

```
db.movies.find({ cast: "Charles Kayser" })
```

13. Retrieve all movies with title, languages, released, directors, writers, countries from the 'movies' collection in MongoDB that released on May 9, 1893.

```
db.movies.find({

 released: {

 $date: ISODate("1893-05-09")

 }

})
```

14. Retrieve all movies with title, languages, released, directors, writers, countries from the 'movies' collection in MongoDB that have a word "scene" in

```
ttitldb.movies.find({ title: { $regex: /scene/i } }) // $regex: regular expression, i: case-insensitive
```



| <b>Evaluation<br/>Procedure</b>  | <b>Marks awarded</b> |
|----------------------------------|----------------------|
| <b>PL/SQL<br/>Procedure(5)</b>   |                      |
| <b>Program/Execution<br/>(5)</b> |                      |
| <b>Viva(5)</b>                   |                      |
| <b>Total (15)</b>                |                      |
| <b>Faculty Signature</b>         |                      |